

ZNEUS - Project 1

```
In [8]: !pip install shapely
```

```
Requirement already satisfied: shapely in c:\users\matze\appdata\local\programs\python\python312\lib\site-packages (2.1.2)
Requirement already satisfied: numpy>=1.21 in c:\users\matze\appdata\local\programs\python\python312\lib\site-packages (from shapely) (2.2.6)
[notice] A new release of pip is available: 24.2 -> 25.3
[notice] To update, run: python.exe -m pip install --upgrade pip
```

```
In [9]: import pandas as pd
```

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
import pylab as py
from scipy import stats
from scipy.stats import pearsonr
import sklearn
import sklearn.preprocessing as preprocessing
from sklearn.feature_selection import SelectKBest, f_regression
from sklearn.pipeline import Pipeline
from shapely import Polygon, Point
from typing import List
import os
import joblib
import plotly.io as pio
TRANSFORMED_PATH = "data/transformed"
r_seed = 42
```

1. EDA

1.1 Loading Dataset

```
In [10]: df = pd.read_csv('data/houses.csv')
df.head()
```

```
Out[10]:   median_house_value  median_income  housing_median_age  total_rooms  total_bedrooms
0          452600.0         8.3252           41.0            880.0            129.0
1          358500.0         8.3014           21.0            7099.0           1106.0
2          352100.0         7.2574           52.0            1467.0            190.0
3          341300.0         5.6431           52.0            1274.0            235.0
4          342200.0         3.8462           52.0            1627.0            280.0
```

1.2 Info about dataset

```
In [11]: print(df.info(),"\n")

print("Unique values:")
for col in df.columns:
    print("\t",col + " počet unikátnych záznamov: ", len(df[col].unique()))

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   median_house_value 20640 non-null   float64
 1   median_income       20640 non-null   float64
 2   housing_median_age 20640 non-null   float64
 3   total_rooms         20640 non-null   float64
 4   total_bedrooms      20640 non-null   float64
 5   population          20640 non-null   float64
 6   households          20640 non-null   float64
 7   latitude             20640 non-null   float64
 8   longitude            20640 non-null   float64
dtypes: float64(9)
memory usage: 1.4 MB
None

Unique values:
    median_house_value počet unikátnych záznamov: 3842
    median_income       počet unikátnych záznamov: 12928
    housing_median_age  počet unikátnych záznamov: 52
    total_rooms          počet unikátnych záznamov: 5926
    total_bedrooms        počet unikátnych záznamov: 1928
    population           počet unikátnych záznamov: 3888
    households           počet unikátnych záznamov: 1815
    latitude              počet unikátnych záznamov: 862
    longitude             počet unikátnych záznamov: 844
```

This dataset contains 9 columns and 20640 entries, all columns are of type float64, therefore are numerical.

Some columns have a lot of unique values.

Target column is **median_house_value**.

1.3 Duplicates

```
In [12]: duplicates = df.duplicated()
print(df[duplicates])

Empty DataFrame
Columns: [median_house_value, median_income, housing_median_age, total_rooms, total_
bedrooms, population, households, latitude, longitude]
Index: []
```

This implicates, that there are no duplicates.

1.4 Missing data

```
In [13]: print(df.isna().sum())
```

```
median_house_value      0
median_income          0
housing_median_age     0
total_rooms            0
total_bedrooms         0
population             0
households             0
latitude               0
longitude              0
dtype: int64
```

There are also no missing data, what a nice dataset.

1.5 Outliers

```
In [14]: def count_outliers(column):
    Q1 = column.quantile(0.25)
    Q3 = column.quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR
    return column[(column < lower) | (column > upper)].shape[0]
```

```
#apply to all columns
outliers_count = df.apply(count_outliers)

print(outliers_count)
```

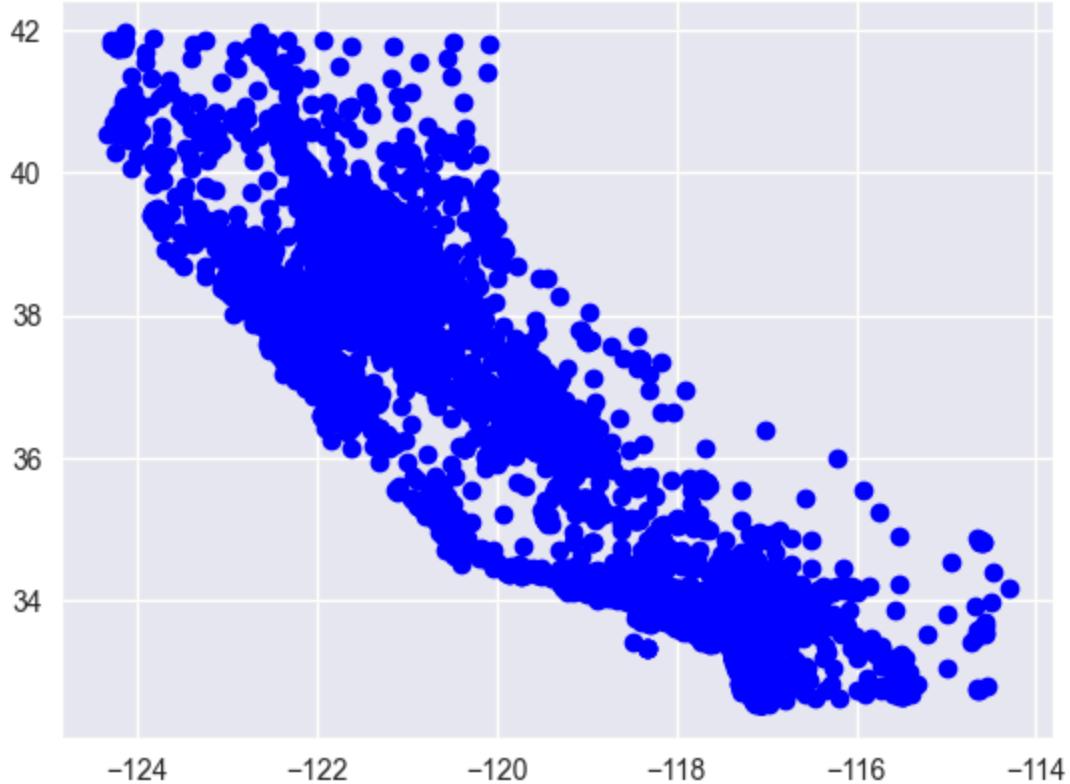
```
median_house_value      1071
median_income           681
housing_median_age      0
total_rooms             1287
total_bedrooms          1282
population              1196
households              1220
latitude                0
longitude               0
dtype: int64
```

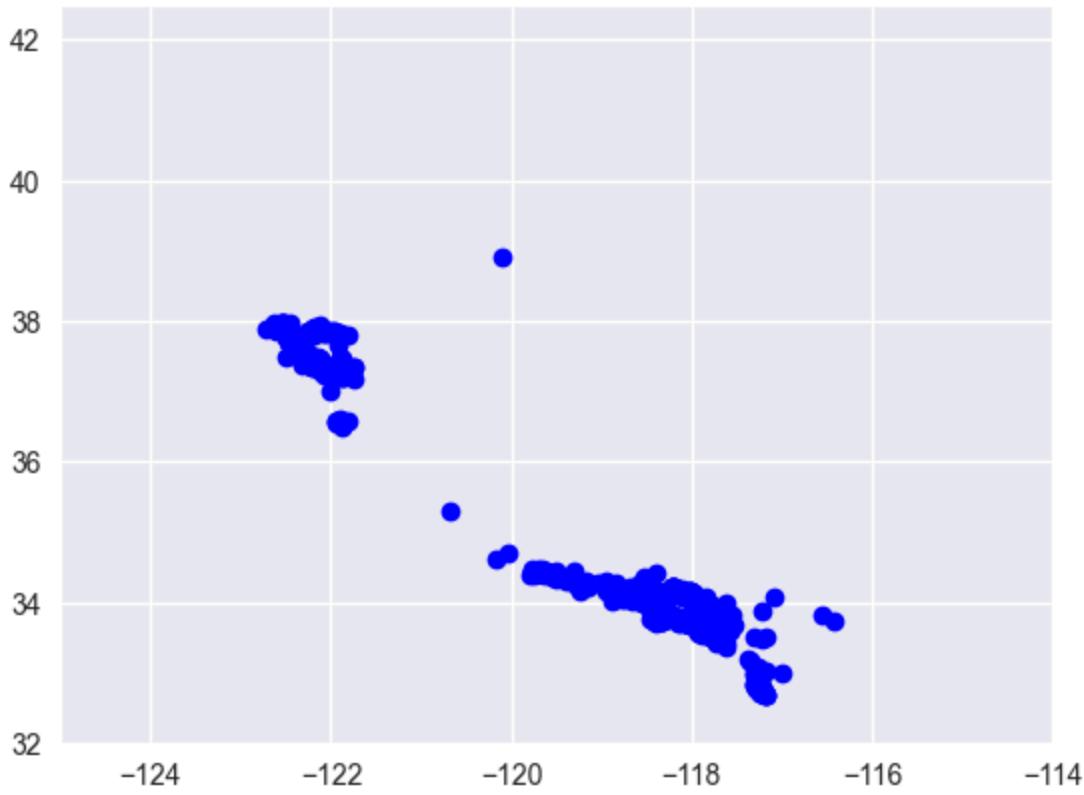
We can see, that we have numerous outliers in some columns.

```
In [15]: def get_outliers(column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR
```

```
return df[(df[column] < lower) | (df[column] > upper)]\n\nouts = get_outliers("median_house_value")\nhouse_outs = get_outliers("median_income")\nplt.figure()\nplt.scatter(df["longitude"], df["latitude"], color='blue', marker='o')\nplt.figure()\nplt.scatter(outs["longitude"], outs["latitude"], color='blue', marker='o')\nplt.xlim(-125, -114)\nplt.ylim(32, 42.5)
```

Out[15]: (32.0, 42.5)



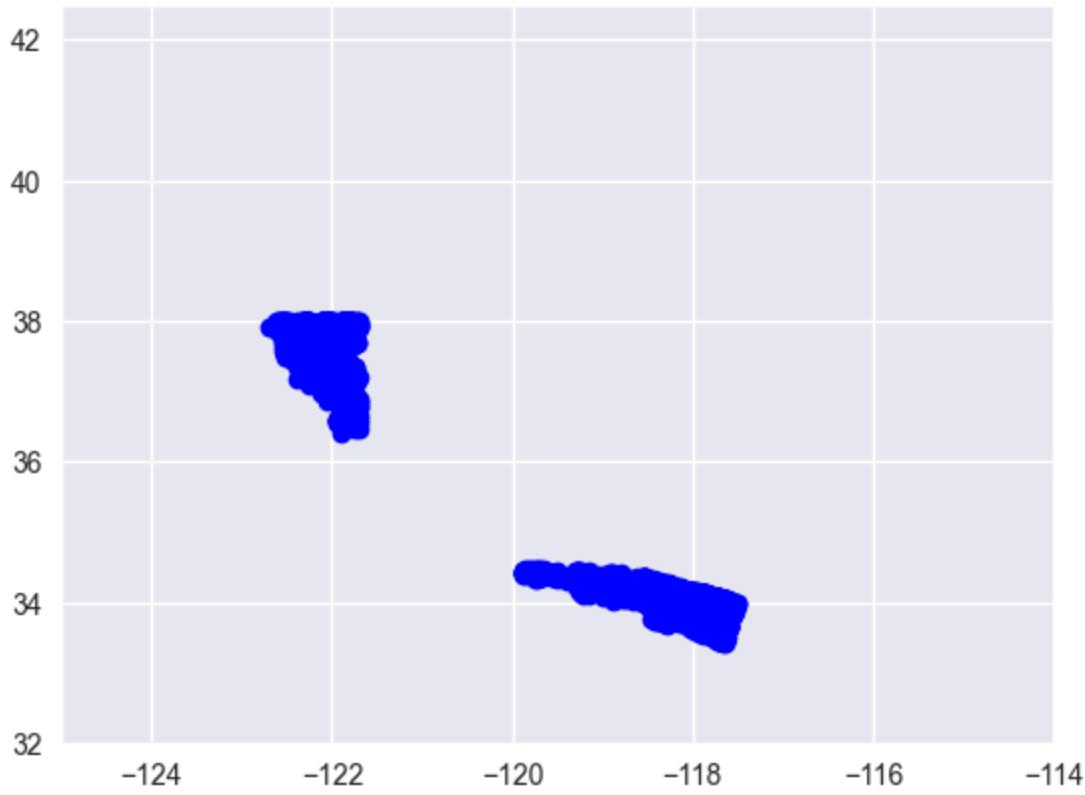


```
In [16]: def set_is_san_francisco(df: pd.DataFrame):
    df["is_san_francisco"] = ((df["longitude"] <= -121.7) & (df["longitude"] > -122
                                & (df["latitude"] > 36.3) & (df["latitude"] <= 38)).astype(int)

    la_polygon = Polygon([
        (-119.9, 34.5),
        (-118.6, 34.4),
        (-117.5, 34),
        (-117.65, 33.4),
        (-119.9, 34.4),
    ])
    def set_is_los_angeles(df: pd.DataFrame):
        df["is_los_angeles"] = df.apply(
            lambda row: int(la_polygon.contains(Point(row["longitude"], row["latitude"]))
                           axis=1
            )

    set_is_san_francisco(df)
    set_is_los_angeles(df)
    plt.figure()
    _df = df[(df["is_san_francisco"] == 1) | (df["is_los_angeles"] == 1)]
    plt.scatter(_df["longitude"], _df["latitude"], color='blue', marker='o')
    plt.xlim(-125, -114)
    plt.ylim(32, 42.5)
```

Out[16]: (32.0, 42.5)



As we can see most of the outliers are focused in specific areas indicating that they are probably result of being located in big cities such as LA

1.5.1 Outlier imputation

```
In [17]: def clip_outliers(column):
    Q1 = column.quantile(0.25)
    Q3 = column.quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR

    return column.clip(lower, upper) #if value is Lower then min, replace it with m

#apply to all
df.apply(clip_outliers)
```

Out[17]:

	median_house_value	median_income	housing_median_age	total_rooms	total_bedrooms
0	452600.0	8.013025	41.0	880.000	1
1	358500.0	8.013025	21.0	5698.375	1
2	352100.0	7.257400	52.0	1467.000	1
3	341300.0	5.643100	52.0	1274.000	2
4	342200.0	3.846200	52.0	1627.000	2
...
20635	78100.0	1.560300	25.0	1665.000	1
20636	77100.0	2.556800	18.0	697.000	1
20637	92300.0	1.700000	17.0	2254.000	1
20638	84700.0	1.867200	18.0	1860.000	1
20639	89400.0	2.388600	16.0	2785.000	1

20640 rows × 11 columns

1.6 Valid data ranges

In [18]:

```
#Latitude
invalid_lat = df[(df['latitude'] < -90) | (df['latitude'] > 90)]
print("Počet invalid latitude:", invalid_lat.shape[0])

#Longitude
invalid_long = df[(df['longitude'] < -180) | (df['longitude'] > 180)]
print("Počet invalid longitude:", invalid_long.shape[0])

#negative values in other columns
df.min()
```

Počet invalid latitude: 0

Počet invalid longitude: 0

Out[18]:

median_house_value	14999.0000
median_income	0.4999
housing_median_age	1.0000
total_rooms	2.0000
total_bedrooms	1.0000
population	3.0000
households	1.0000
latitude	32.5400
longitude	-124.3500
is_san_francisco	0.0000
is_los_angeles	0.0000
dtype:	float64

There are no negative values or zero values in other columns, this is okay. Lat and Long are both in valid ranges.

1.7 Summary statistics

```
In [19]: df.describe()
```

	median_house_value	median_income	housing_median_age	total_rooms	total_bedrooms
count	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000
mean	206855.816909	3.870671	28.639486	2635.763081	537.895000
std	115395.615874	1.899822	12.585558	2181.615252	421.242000
min	14999.000000	0.499900	1.000000	2.000000	1.000000
25%	119600.000000	2.563400	18.000000	1447.750000	295.000000
50%	179700.000000	3.534800	29.000000	2127.000000	435.000000
75%	264725.000000	4.743250	37.000000	3148.000000	647.000000
max	500001.000000	15.000100	52.000000	39320.000000	6445.000000

```
In [20]: print("Other stats:")
for col in df.columns:
    print("\t", col, ":")
    print("\t\t", "mean", np.mean(df[col]))
    print("\t\t", "median", np.median(df[col]))
    print("\t\t", "mode", stats.mode(df[col])[0])
    print("\t\t", "variance", np.var(df[col]), "\n")
```

Other stats:

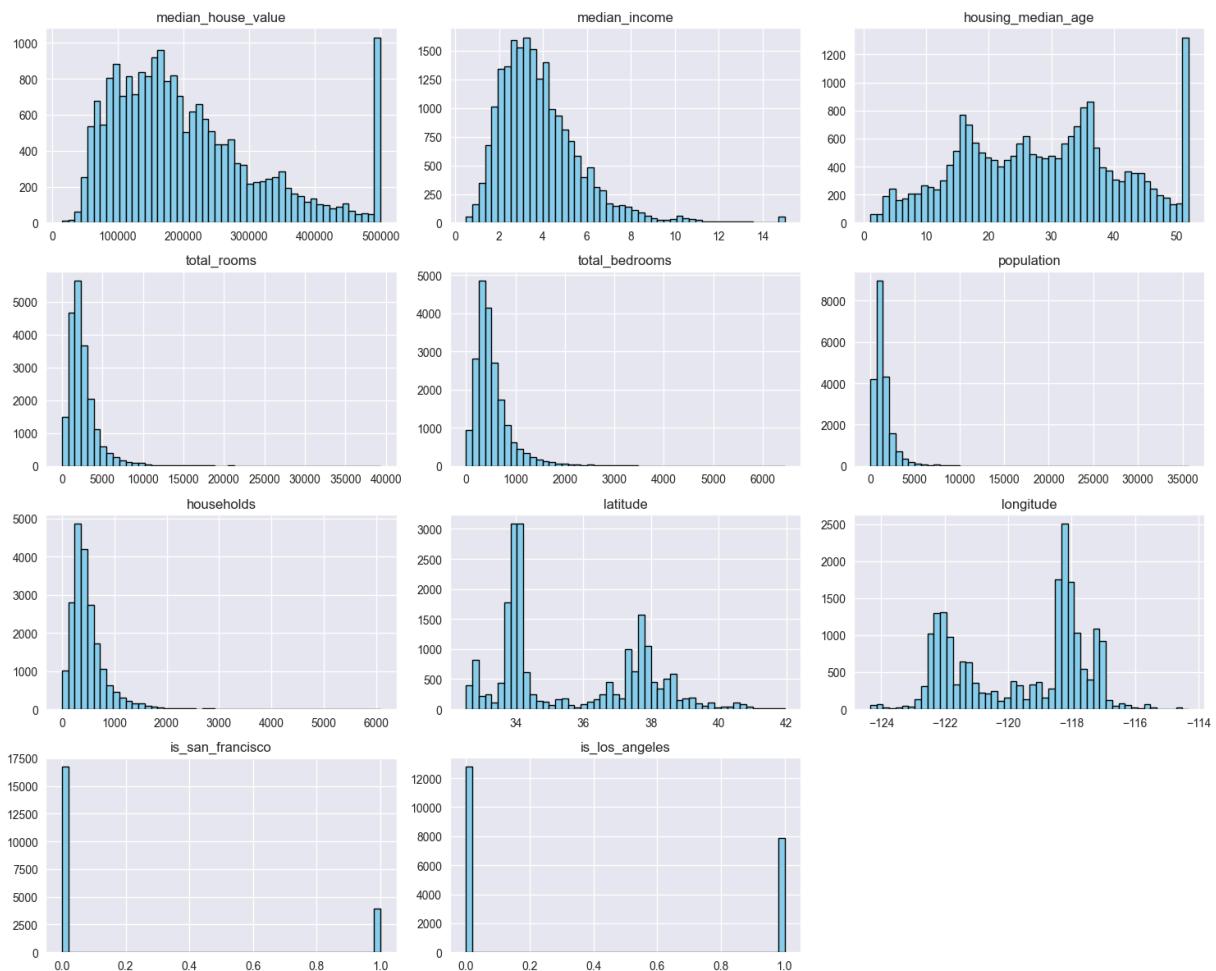
```
median_house_value :  
    mean 206855.81690891474  
    median 179700.0  
    mode 500001.0  
    variance 13315503000.818077  
  
median_income :  
    mean 3.8706710029069766  
    median 3.5347999999999997  
    mode 3.125  
    variance 3.609147689697444  
  
housing_median_age :  
    mean 28.639486434108527  
    median 29.0  
    mode 52.0  
    variance 158.38858617035862  
  
total_rooms :  
    mean 2635.7630813953488  
    median 2127.0  
    mode 1527.0  
    variance 4759214.512668024  
  
total_bedrooms :  
    mean 537.8980135658915  
    median 435.0  
    mode 280.0  
    variance 177441.20088752697  
  
population :  
    mean 1425.4767441860465  
    median 1166.0  
    mode 891.0  
    variance 1282408.3220366866  
  
households :  
    mean 499.5396802325581  
    median 409.0  
    mode 306.0  
    variance 146168.95772780472  
  
latitude :  
    mean 35.63186143410853  
    median 34.26  
    mode 34.06  
    variance 4.562071602892517  
  
longitude :  
    mean -119.56970445736432  
    median -118.49  
    mode -118.31  
    variance 4.0139448835847835  
  
is_san_francisco :
```

```
mean 0.19050387596899224
median 0.0
mode 0
variance 0.15421214920978307
```

```
is_los_angeles :
mean 0.3814437984496124
median 0.0
mode 0
variance 0.2359444270739439
```

1.8 Visualize data

```
In [21]: df.hist(bins=50, figsize=(15, 12), color='skyblue', edgecolor='black')
plt.tight_layout()
plt.show()
```



None of these distributions appear normal, lets run some tests.

1.8.1 Test for normality

```
In [22]: def check_normal(col):
print(col.name,":")
#draw qqplot
```

```

sm.qqplot(col, line='45')
py.show()

#perform KS test, we dont do shapiro because we have more than 5000 samples
kolmogorov_smirnov = stats.kstest(col, "norm")

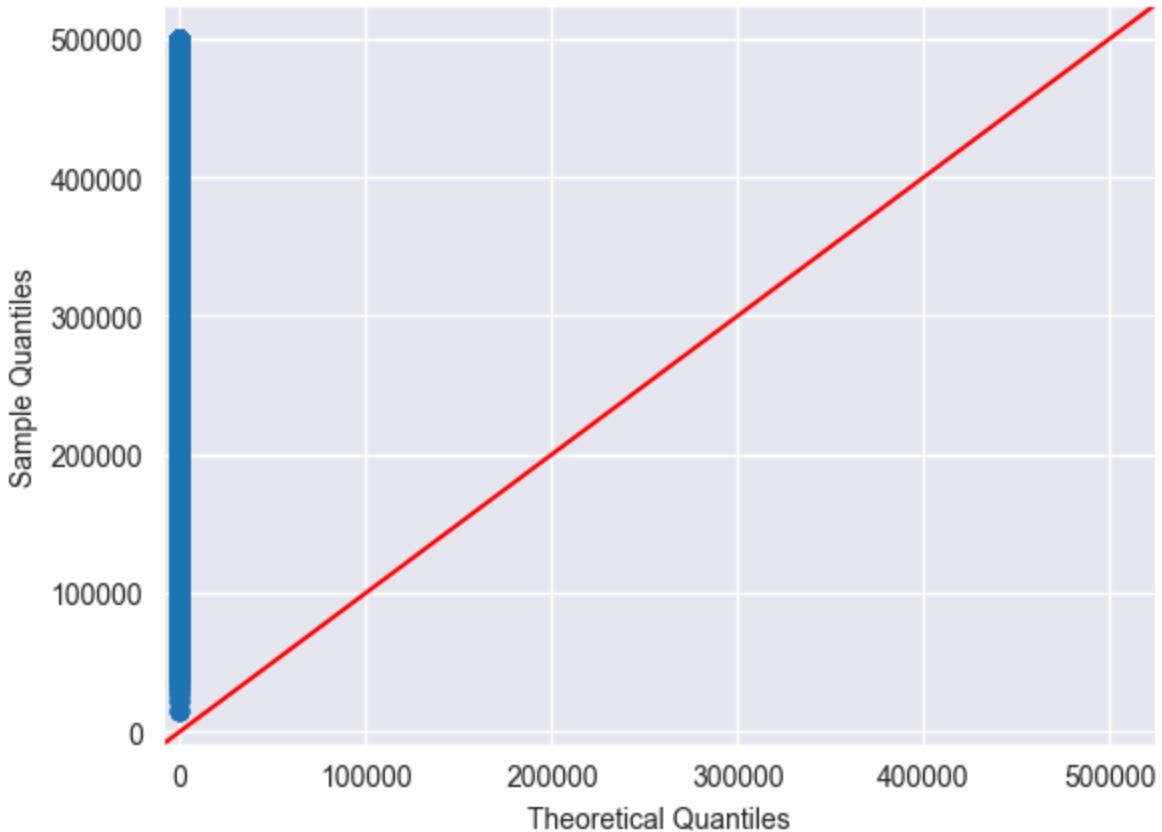
#interpret results
alpha = 0.05

print("Kolmogorov-Smirnov test:", kolmogorov_smirnov)
if kolmogorov_smirnov.pvalue > alpha:
    print('Normal distribution (fail to reject H0)\n')
else:
    print('Another distribution (reject H0)\n')

for col in df:
    check_normal(df[col])

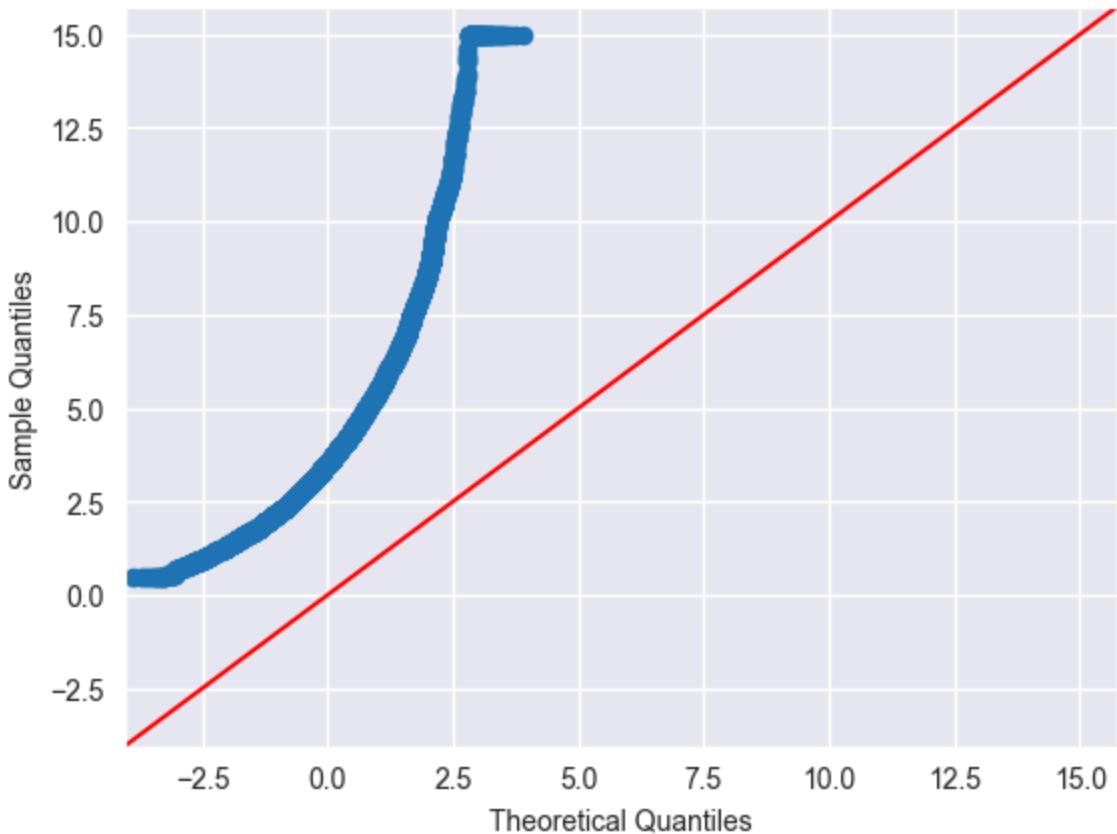
```

median_house_value :



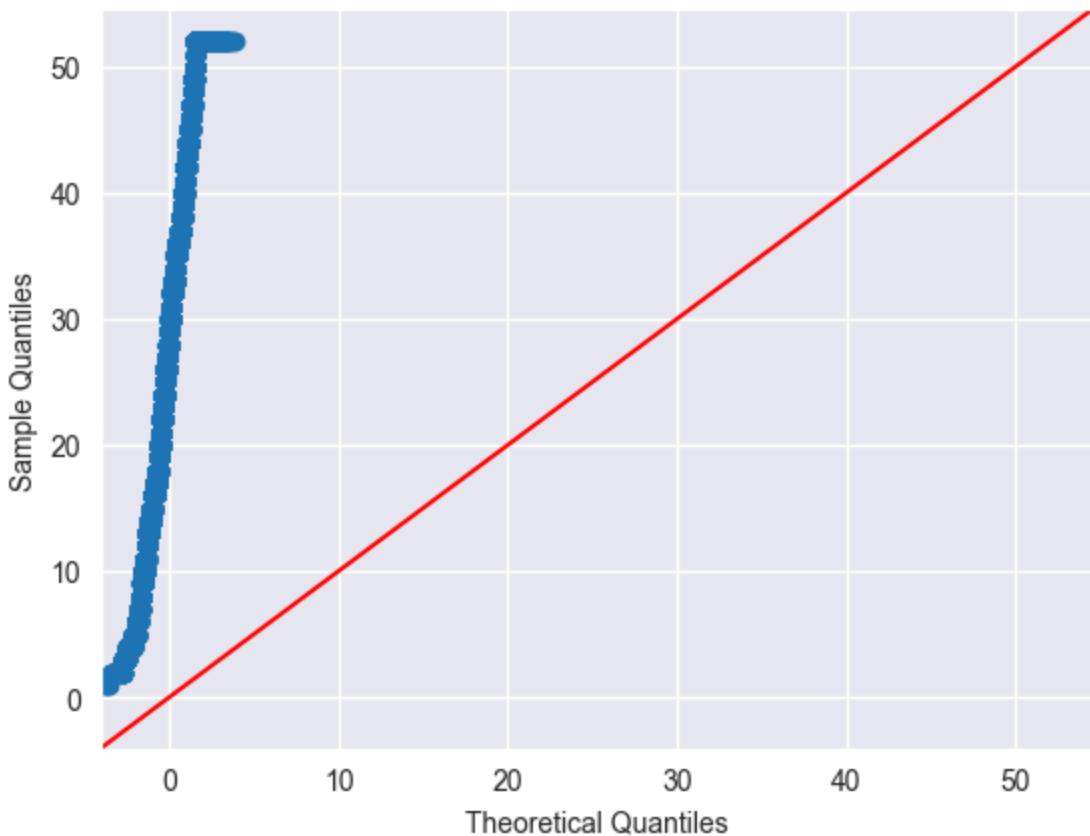
Kolmogorov-Smirnov test: KstestResult(statistic=np.float64(1.0), pvalue=np.float64(0.0), statistic_location=np.float64(14999.0), statistic_sign=np.int8(-1))
Another distribution (reject H0)

median_income :



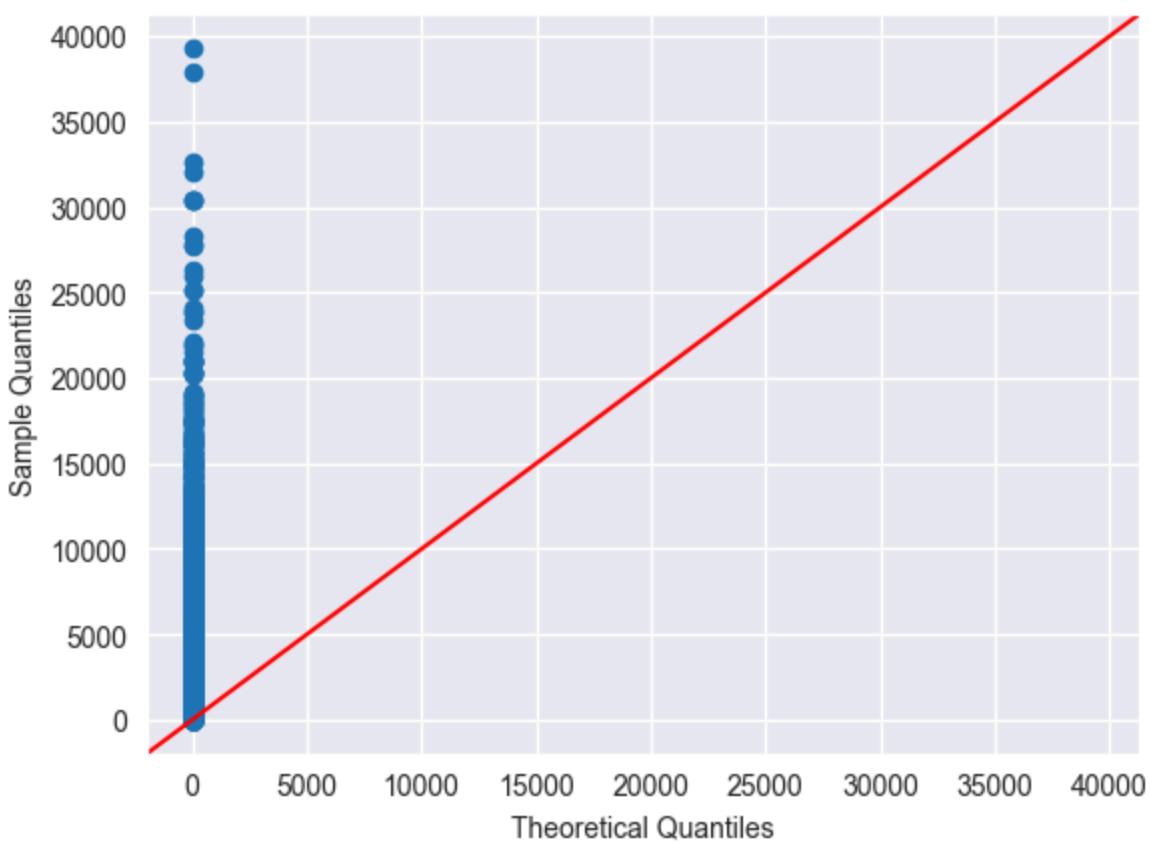
```
Kolmogorov-Smirnov test: KstestResult(statistic=np.float64(0.8953266796146241), pvalue=np.float64(0.0), statistic_location=np.float64(1.5809), statistic_sign=np.int8(-1))  
Another distribution (reject H0)
```

```
housing_median_age :
```



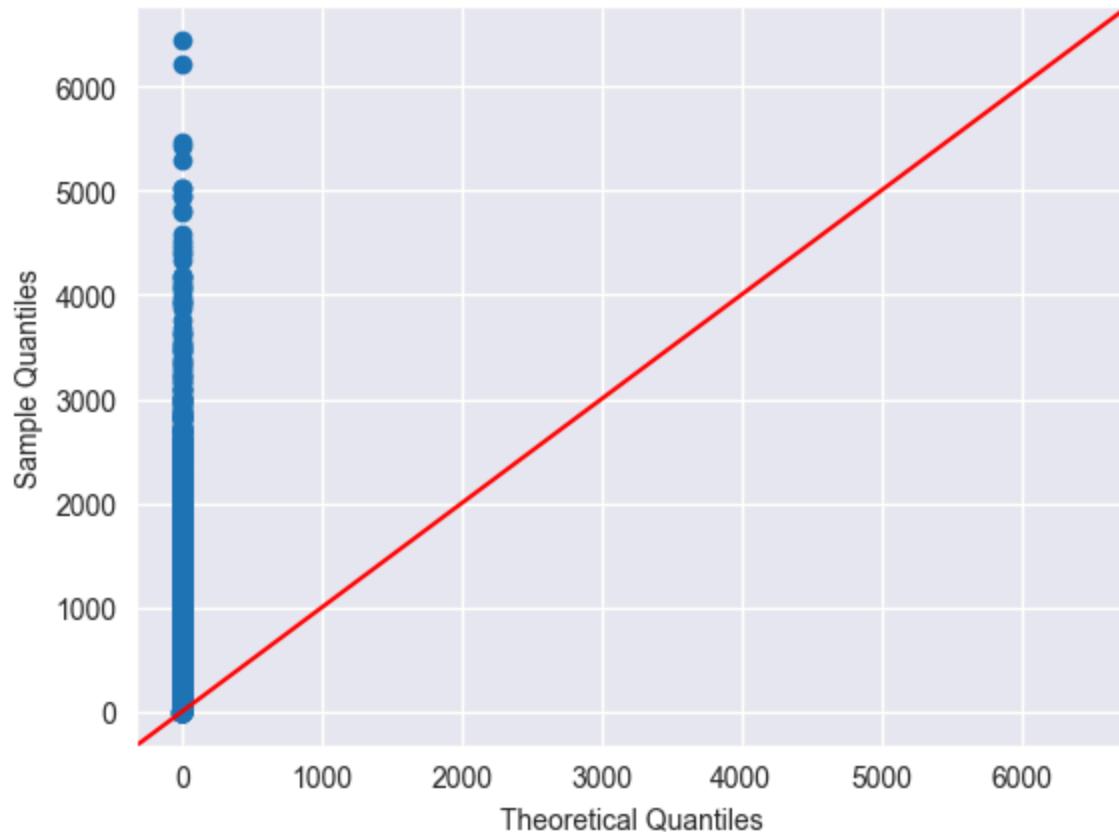
```
Kolmogorov-Smirnov test: KstestResult(statistic=np.float64(0.9956462259993777), pvalue=np.float64(0.0), statistic_location=np.float64(3.0), statistic_sign=np.int8(-1))  
Another distribution (reject H0)
```

total_rooms :



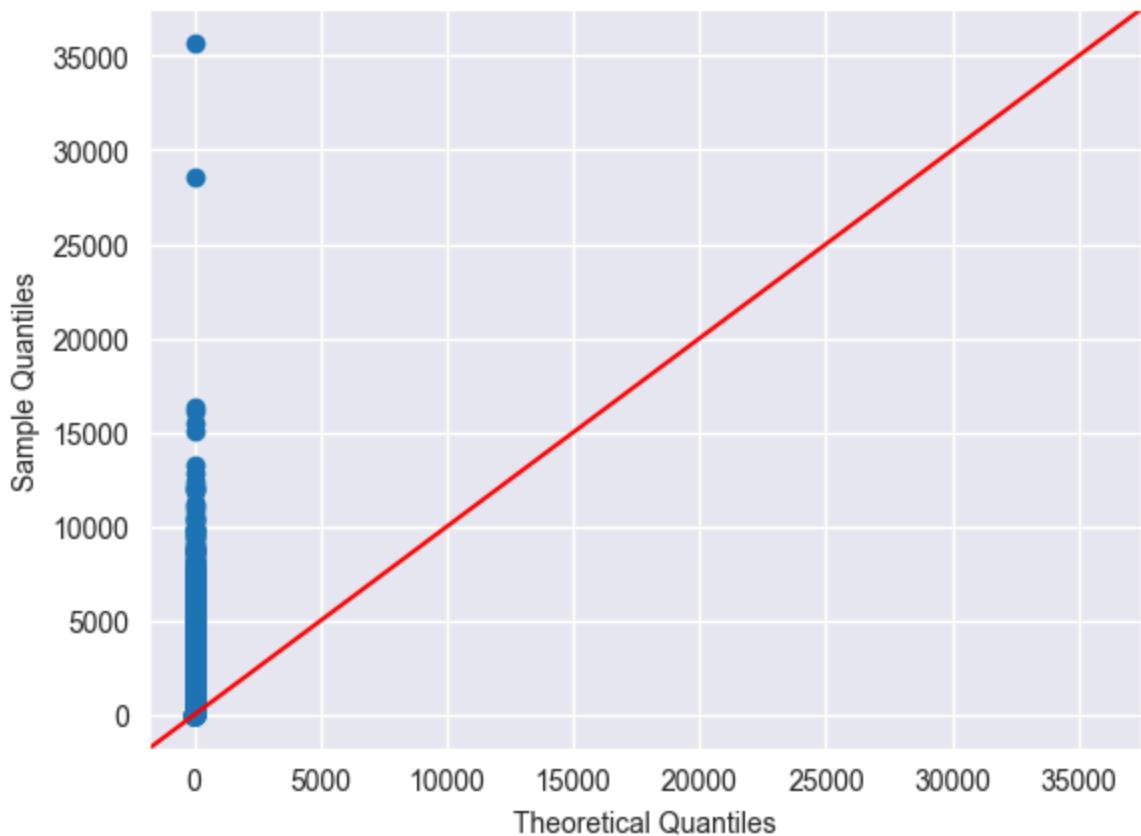
```
Kolmogorov-Smirnov test: KstestResult(statistic=np.float64(0.9999515494010092), pvalue=np.float64(0.0), statistic_location=np.float64(6.0), statistic_sign=np.int8(-1))
Another distribution (reject H0)
```

total_bedrooms :



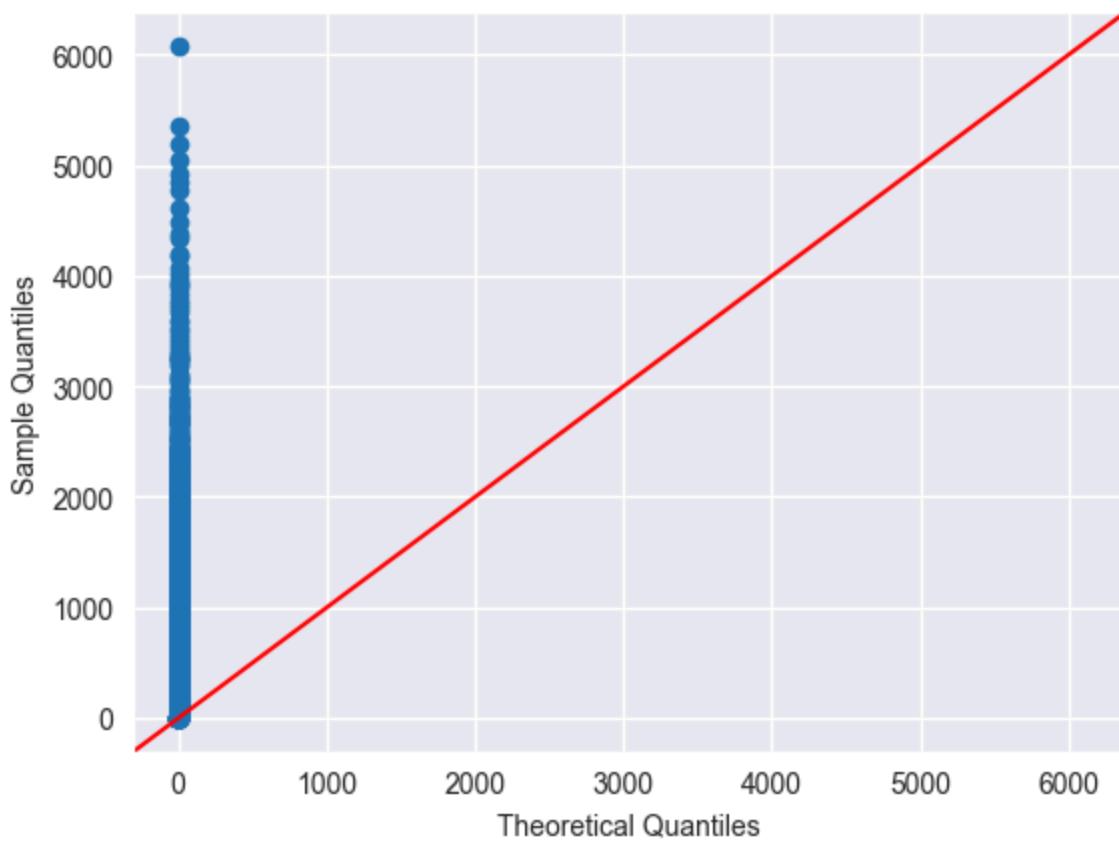
```
Kolmogorov-Smirnov test: KstestResult(statistic=np.float64(0.999580731858942), pvalue=np.float64(0.0), statistic_location=np.float64(4.0), statistic_sign=np.int8(-1))
Another distribution (reject H0)
```

population :



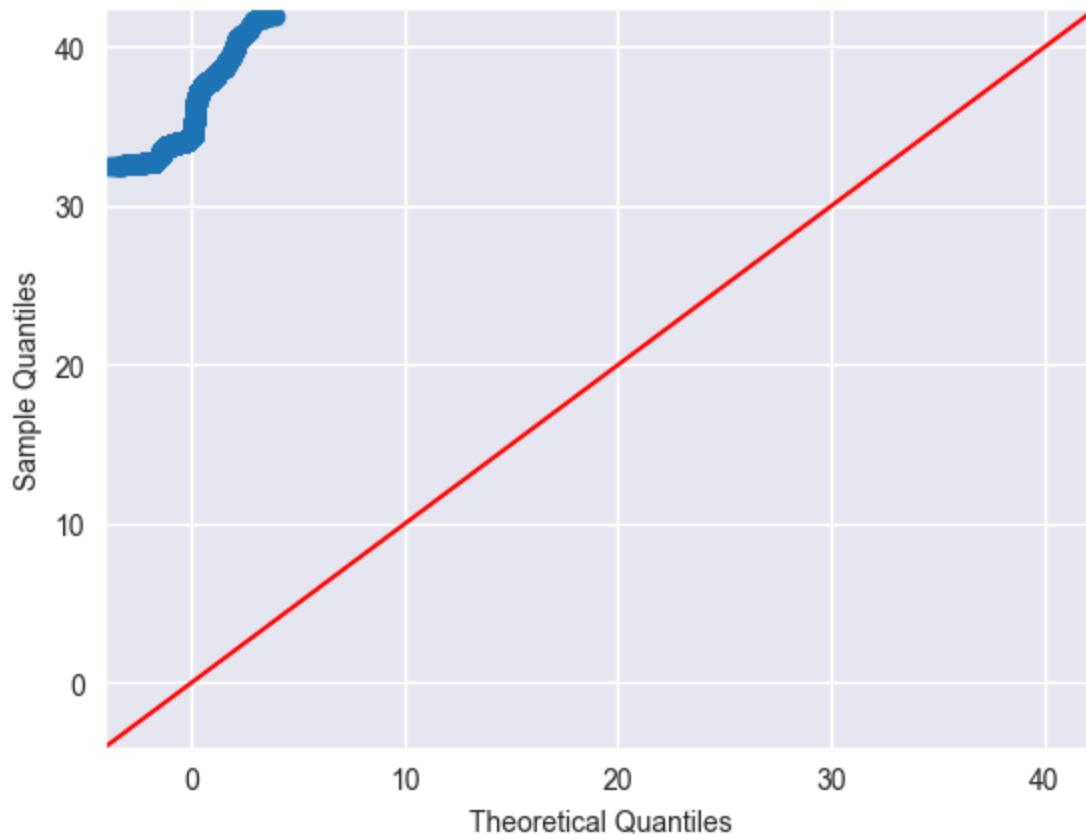
```
Kolmogorov-Smirnov test: KstestResult(statistic=np.float64(0.9999512637360249), pvalue=np.float64(0.0), statistic_location=np.float64(5.0), statistic_sign=np.int8(-1))  
Another distribution (reject H0)
```

households :



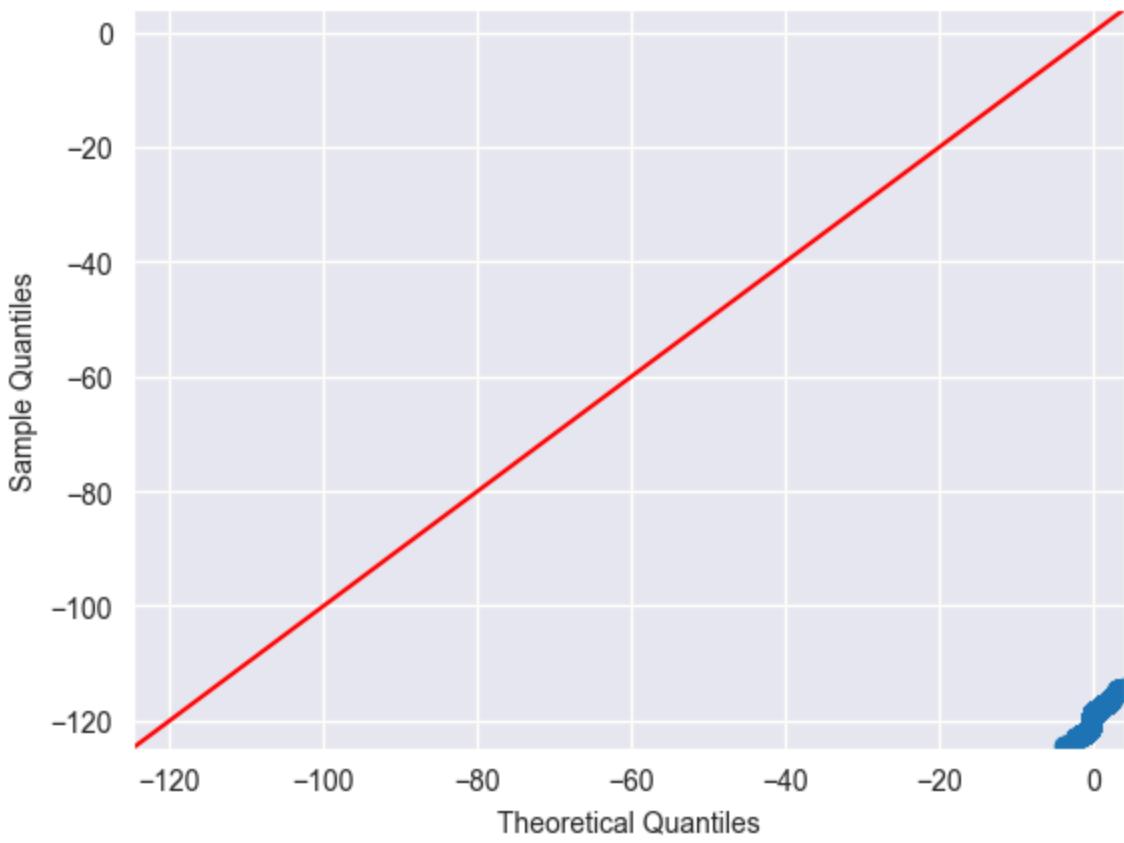
```
Kolmogorov-Smirnov test: KstestResult(statistic=np.float64(0.999580731858942), pvalue=np.float64(0.0), statistic_location=np.float64(4.0), statistic_sign=np.int8(-1))
Another distribution (reject H0)
```

latitude :



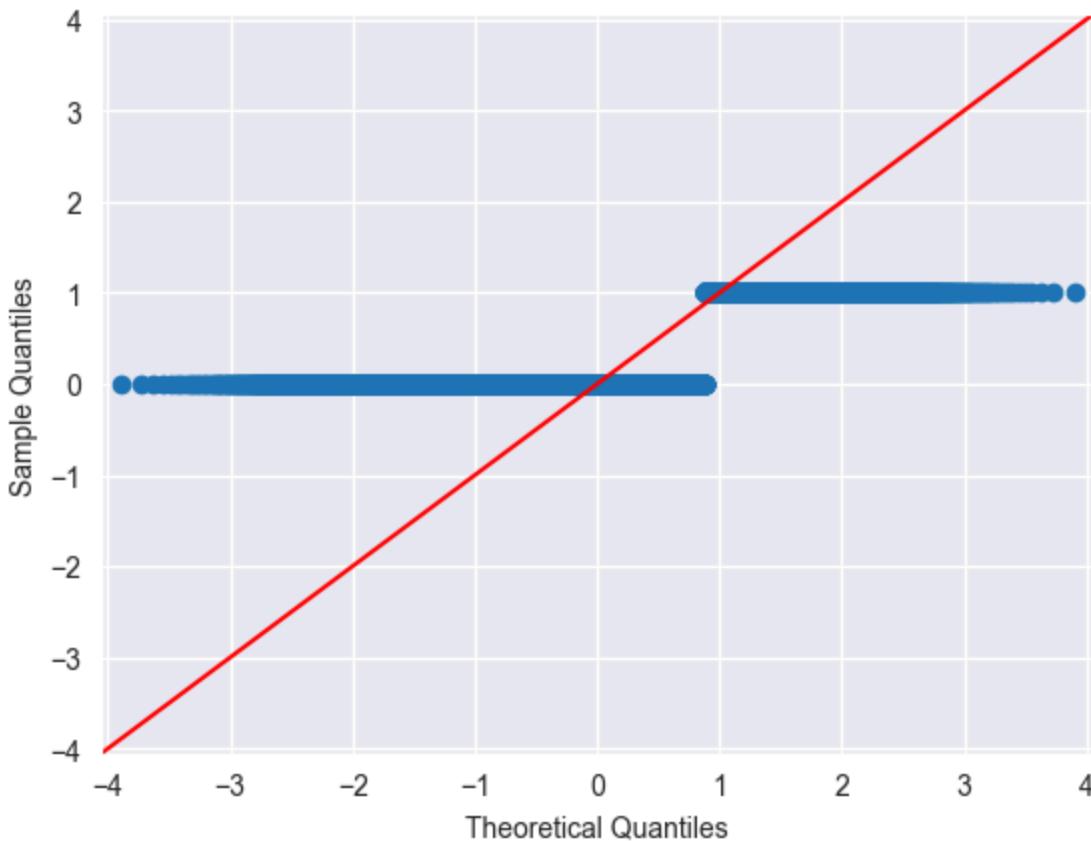
```
Kolmogorov-Smirnov test: KstestResult(statistic=np.float64(1.0), pvalue=np.float64(0.0), statistic_location=np.float64(32.54), statistic_sign=np.int8(-1))
Another distribution (reject H0)
```

longitude :



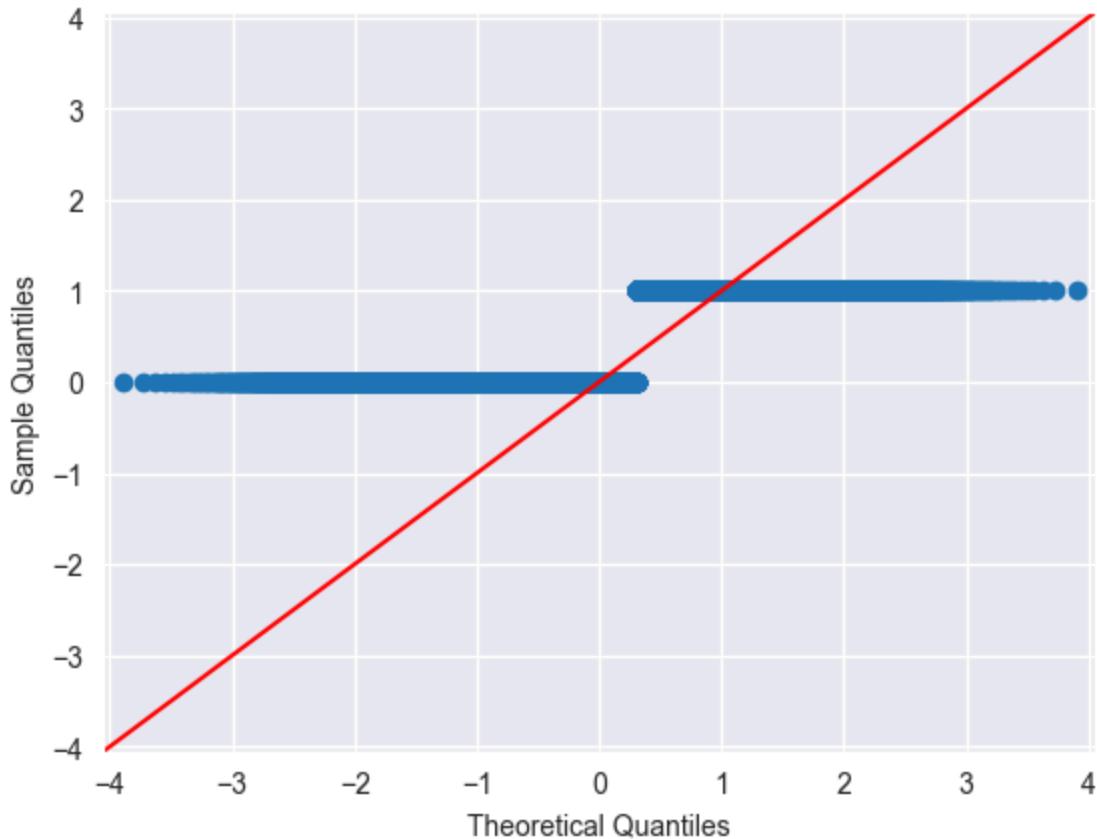
```
Kolmogorov-Smirnov test: KstestResult(statistic=np.float64(1.0), pvalue=np.float64(0.0), statistic_location=np.float64(-114.31), statistic_sign=np.int8(1))  
Another distribution (reject H0)
```

is_san_francisco :



```
Kolmogorov-Smirnov test: KstestResult(statistic=np.float64(0.5), pvalue=np.float64(0.0), statistic_location=np.int64(0), statistic_sign=np.int8(-1))
Another distribution (reject H0)
```

```
is_los_angeles :
```

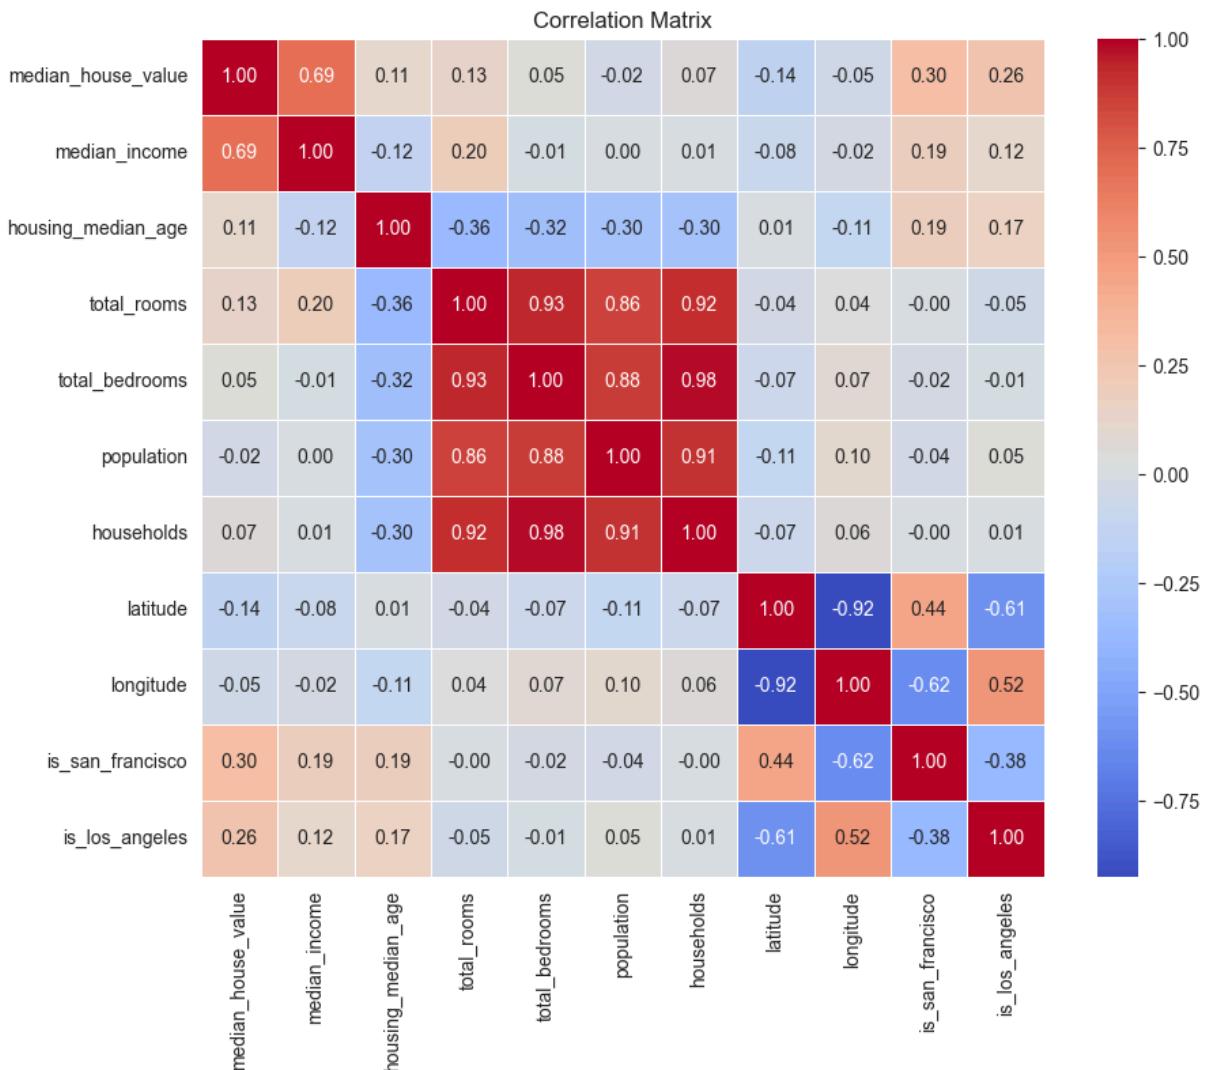


```
Kolmogorov-Smirnov test: KstestResult(statistic=np.float64(0.5), pvalue=np.float64(0.0), statistic_location=np.int64(0), statistic_sign=np.int8(-1))
Another distribution (reject H0)
```

1.9 Correlations

```
In [23]: correlation_matrix = df.corr()

#visualize the correlation matrix in a heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=1)
plt.title('Correlation Matrix')
plt.show()
```



We can see, that columns: median_income (0.69) is strongly positively correlated with the target column.

Other columns are correlated between each other such as: total_rooms - households, population, total_bedrooms, hoouse_median_age and more.

1.10 Hypothesis testing

H0 - median_house_value is not dependent on median_income

H1 - median_house_value increases as median_income increases

1.10.1 Scatterplot and regression plot

```
In [24]: plt.figure(figsize=(8,6))
sns.scatterplot(x='median_income', y='median_house_value', data=df, alpha=0.5)
sns.regplot(x='median_income', y='median_house_value', data=df, scatter=False, color='red')
plt.title('Median House Value vs Median Income')
plt.xlabel('Median Income')
```

```
plt.ylabel('Median House Value')
plt.show()
```



We can see the positive correlation, if we increase median income, we increase median house value as well.

House values are capped at 500000, probably dataset cap. There are notable outliers.

1.10.2 Statistical test

```
In [25]: corr, p_value = pearsonr(df['median_income'], df['median_house_value'])
print(f"Pearson correlation: {corr:.2f}")
print(f"P-value: {p_value:.5f}")
```

Pearson correlation: 0.69

P-value: 0.00000

We can confirm from **p-value < 0.05** that the relation is statistically significant, therefore we reject **H0**.

```
In [26]: stat, p = stats.ks_2samp(df["total_bedrooms"], df["population"])
print(f"KS statistic: {stat}, p-value: {p}")
```

KS statistic: 0.59515503875969, p-value: 0.0

1.11 Evaluation metrics:

- **MSE - Mean Squared Error**, because we are doing Regression task.

Data preprocessing

Dataset division

In [27]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   median_house_value 20640 non-null   float64
 1   median_income       20640 non-null   float64
 2   housing_median_age 20640 non-null   float64
 3   total_rooms         20640 non-null   float64
 4   total_bedrooms      20640 non-null   float64
 5   population          20640 non-null   float64
 6   households          20640 non-null   float64
 7   latitude             20640 non-null   float64
 8   longitude            20640 non-null   float64
 9   is_san_francisco    20640 non-null   int64  
 10  is_los_angeles       20640 non-null   int64  
dtypes: float64(9), int64(2)
memory usage: 1.7 MB
```

In [28]: `# 70 - 15 - 15 split`

```
x = df.drop("median_house_value", axis=1)
y = df[["median_house_value"]]
x_train, x_test, y_train, y_test = sklearn.model_selection.train_test_split(
    x, y, test_size=0.3, random_state=r_seed
)
x_test, x_val, y_test, y_val = sklearn.model_selection.train_test_split(
    x_test, y_test, test_size=0.5, random_state=r_seed
)
x_train.shape, x_test.shape, x_val.shape, y_train.shape, y_test.shape, y_val.shape
```

Out[28]: `((14448, 10), (3096, 10), (3096, 10), (14448, 1), (3096, 1), (3096, 1))`

Normalization

In [29]: `def show_transformations(dfs: List[any], columns: int = 3, bins: int = 50):`

```
n = len(dfs)
rows = (n + columns - 1) // columns

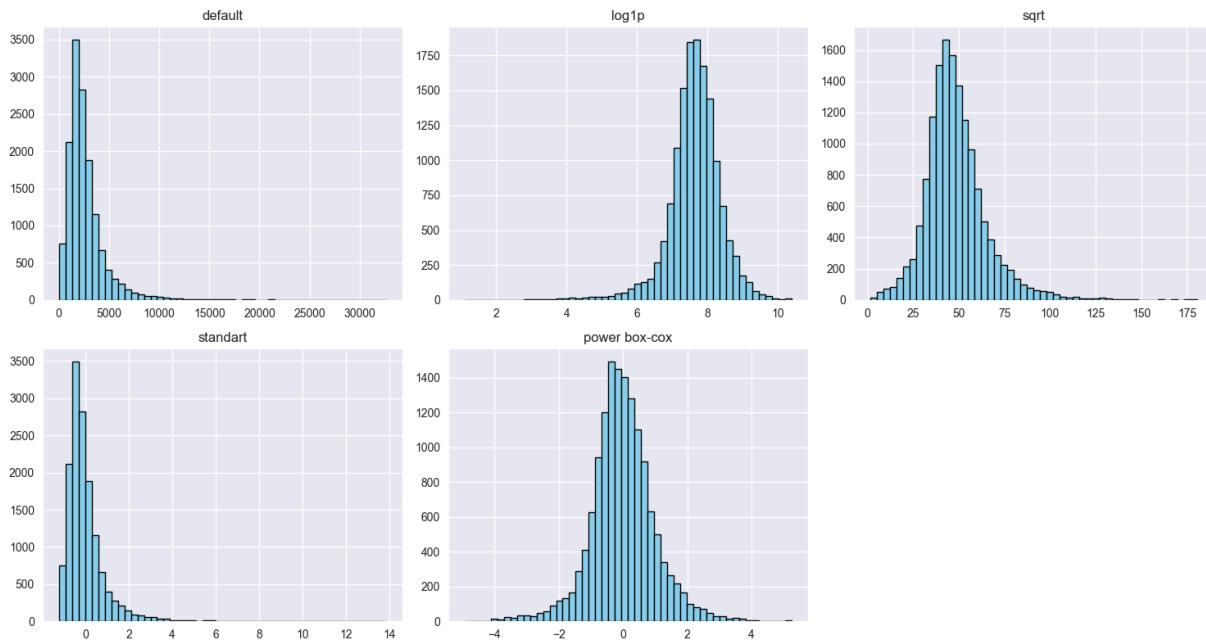
fig, axes = plt.subplots(rows, columns, figsize=(5 * columns, 4 * rows))
axes = axes.flatten()
for i, df in enumerate(dfs):
    axes[i].hist(df[0], bins=bins, color='skyblue', edgecolor='black')
    axes[i].set_title(df[1])
    for j in range(i + 1, len(axes)):
```

```
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()
```

```
In [30]: def test_transformations(data, dfs: List[any] = [], columns: int = 3, bins: int = 5
standart = preprocessing.StandardScaler().fit_transform(data)
power_box = preprocessing.PowerTransformer(method='box-cox').fit_transform(data)
show_transformations([(data, "default"), (np.log1p(data), "log1p"), (np.sqrt(da
(standart, "standart"), (power_box, "power box-cox"), *dfs],
```

total_rooms

```
In [31]: test_transformations(x_train[["total_rooms"]])
```



log transformation skewed data to the left, so it is too much

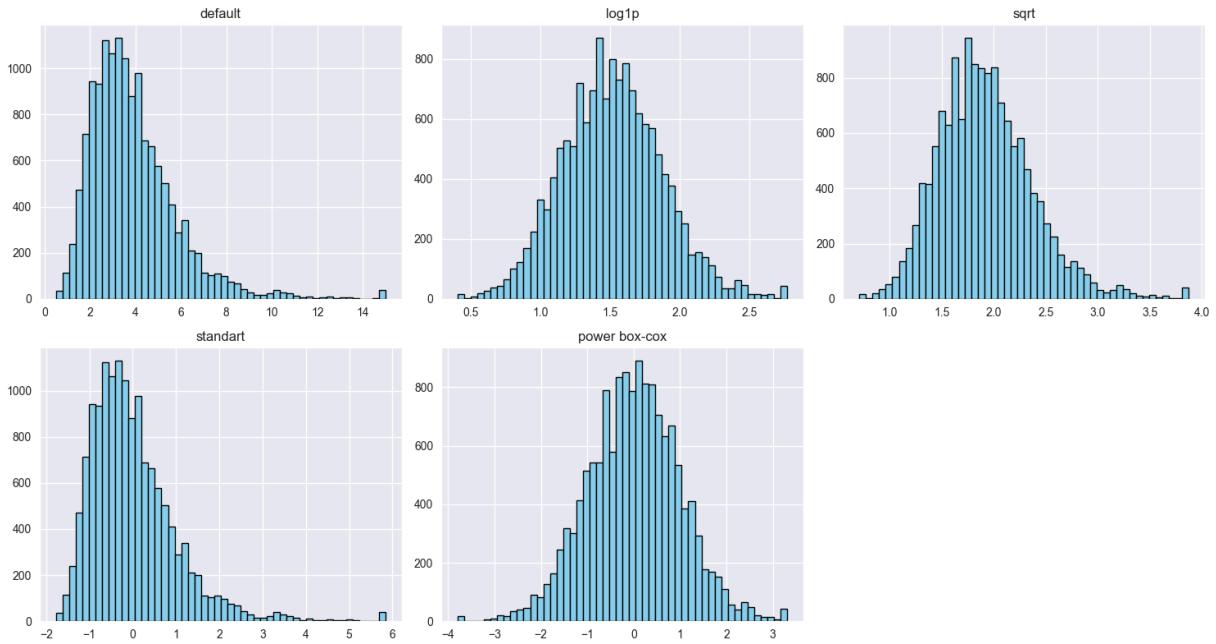
sqrt transformation made the data less right skewed but it is not symmetrical

standard scaler haven't done much that is expected considering that this transformer expects roughly normal distribution

power transformer made the data look roughly symmetrical

median_income

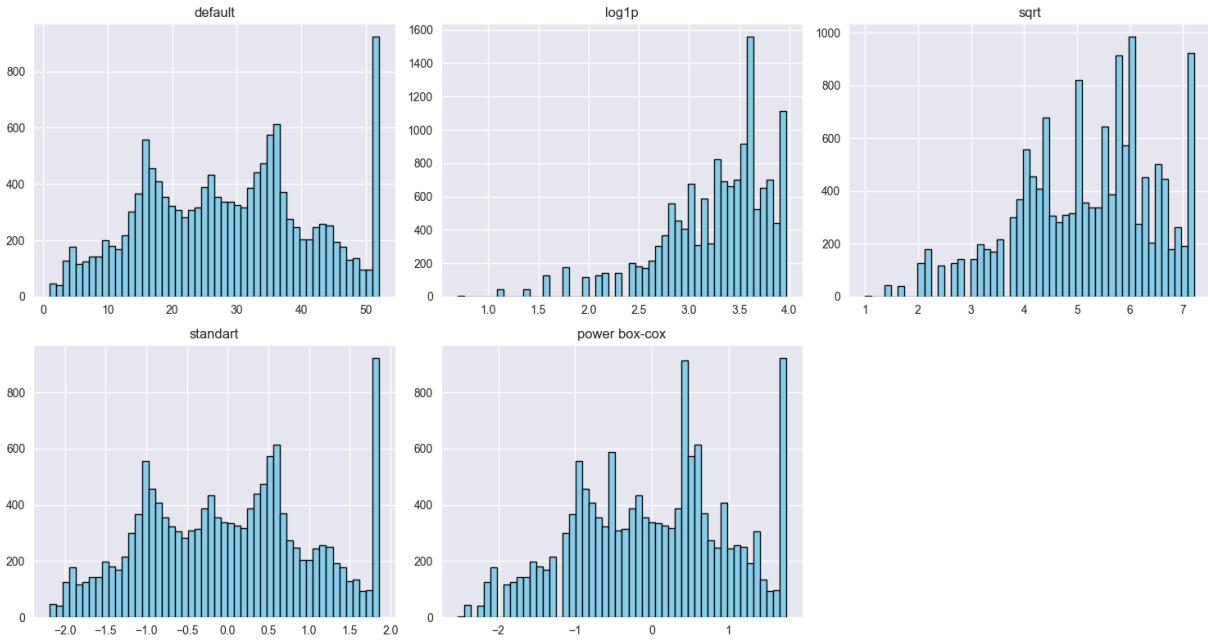
```
In [32]: param = "median_income"
test_transformations(x_train[[param]])
```



Both log and power transformation appear symmetric

housing_median_age

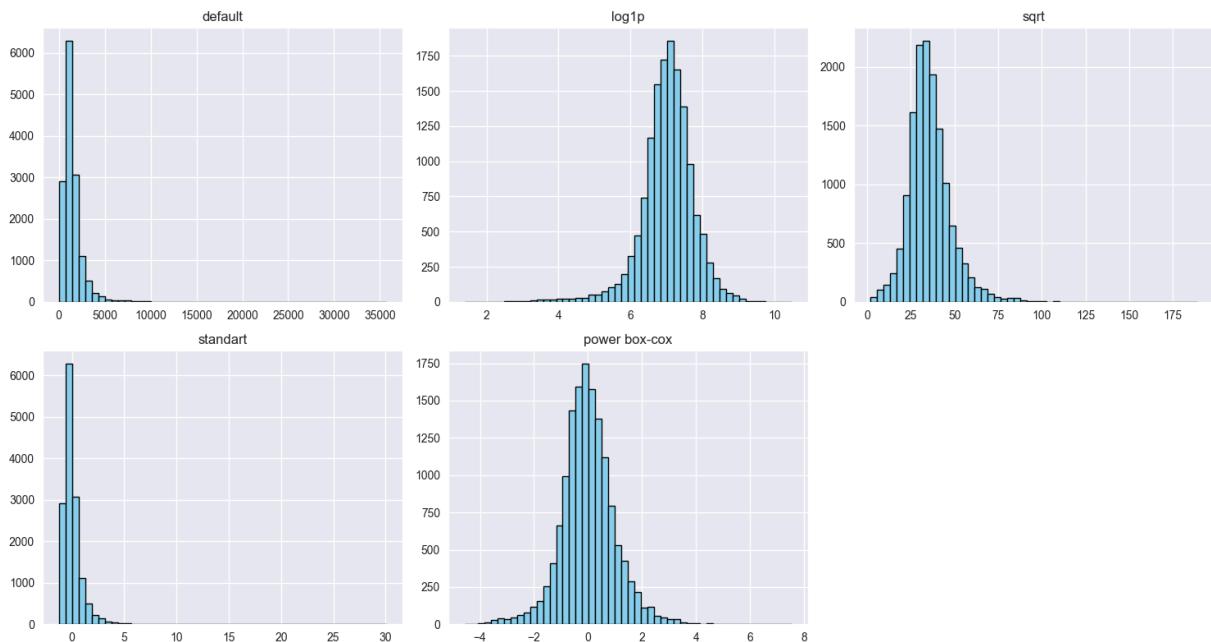
```
In [33]: param = "housing_median_age"
test_transformations(x_train[[param]])
```



None of the transformations above were able to make the data more normal or symmetric, so the Min-Max scaler can be used to preserve the original metric units

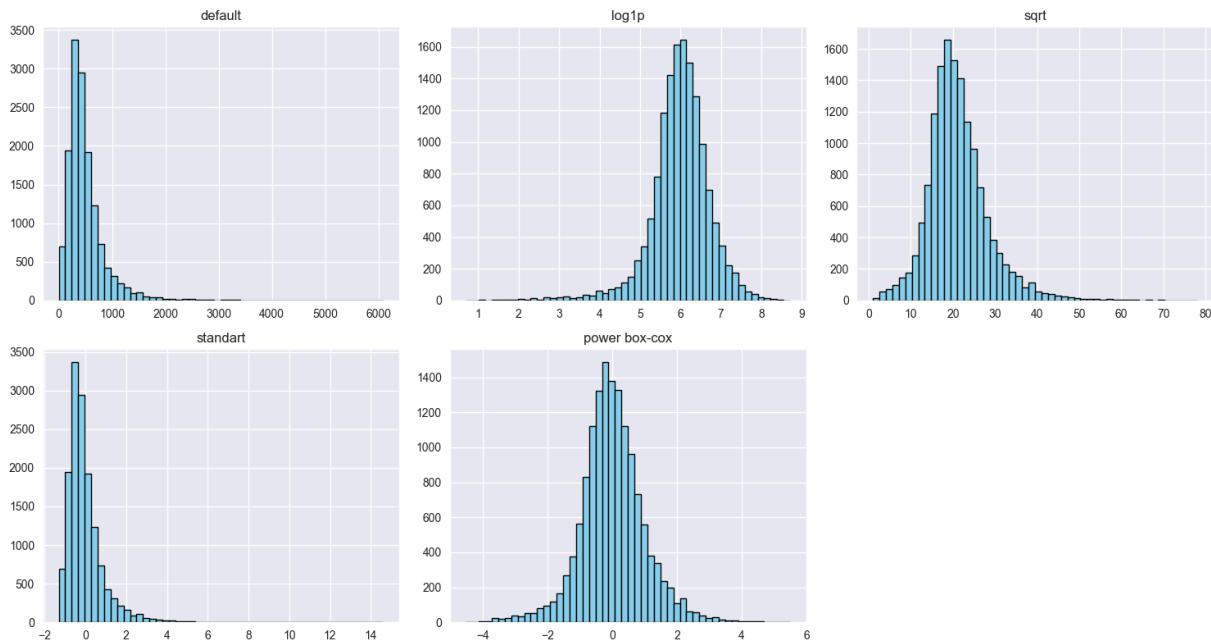
population

```
In [34]: param = "population"
test_transformations(x_train[[param]])
```



households

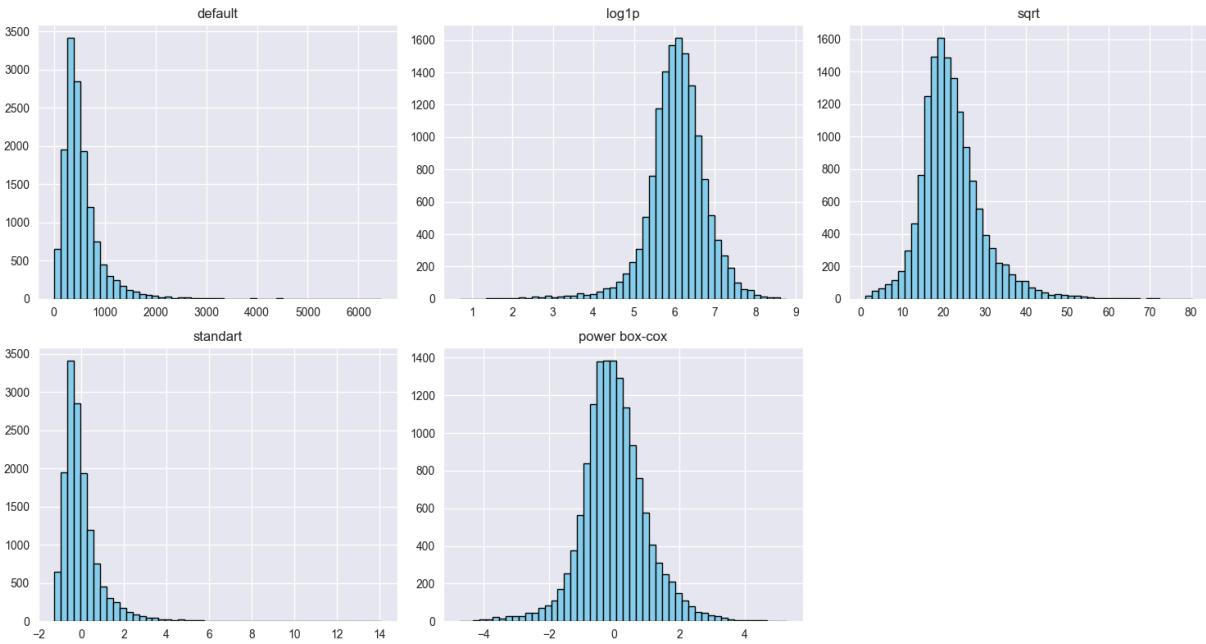
```
In [35]: param = "households"
test_transformations(x_train[[param]])
```



Households appear the least skewed after the power transformation

total_bedrooms

```
In [36]: param = "total_bedrooms"
test_transformations(x_train[[param]])
```



latitude & longitude

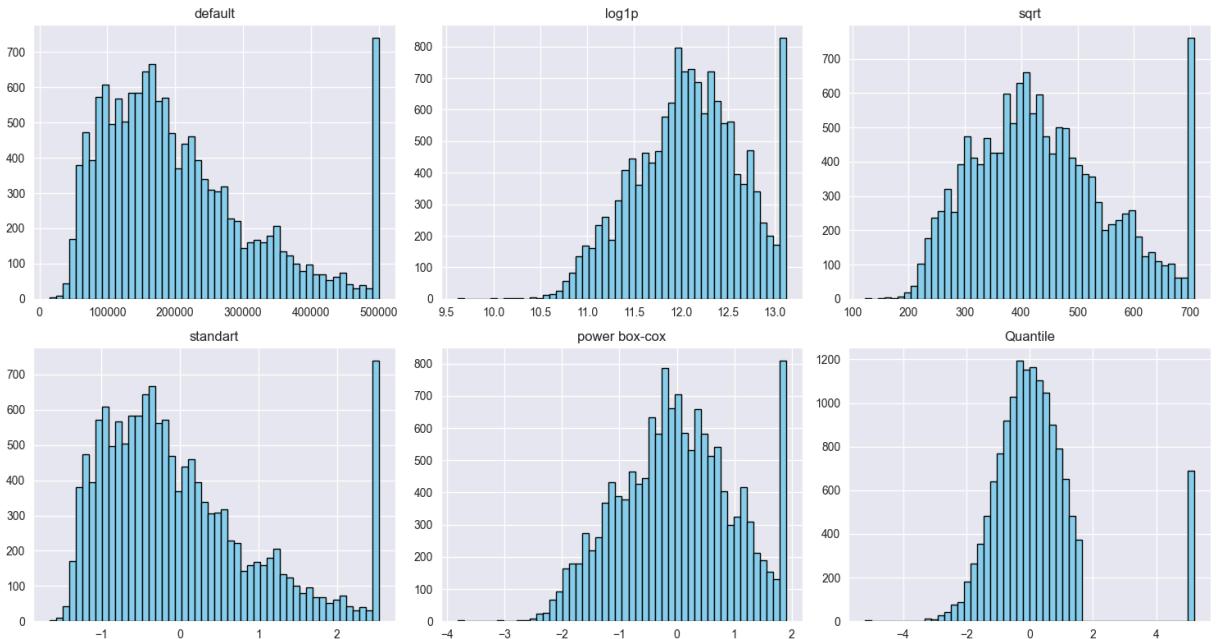
Since it is a location metric and its original units are important MinMax transformer will be used

is_san_francisco & is_los_angeles

These are boolean features so there is no need to normalize or rescale them

median_house_value

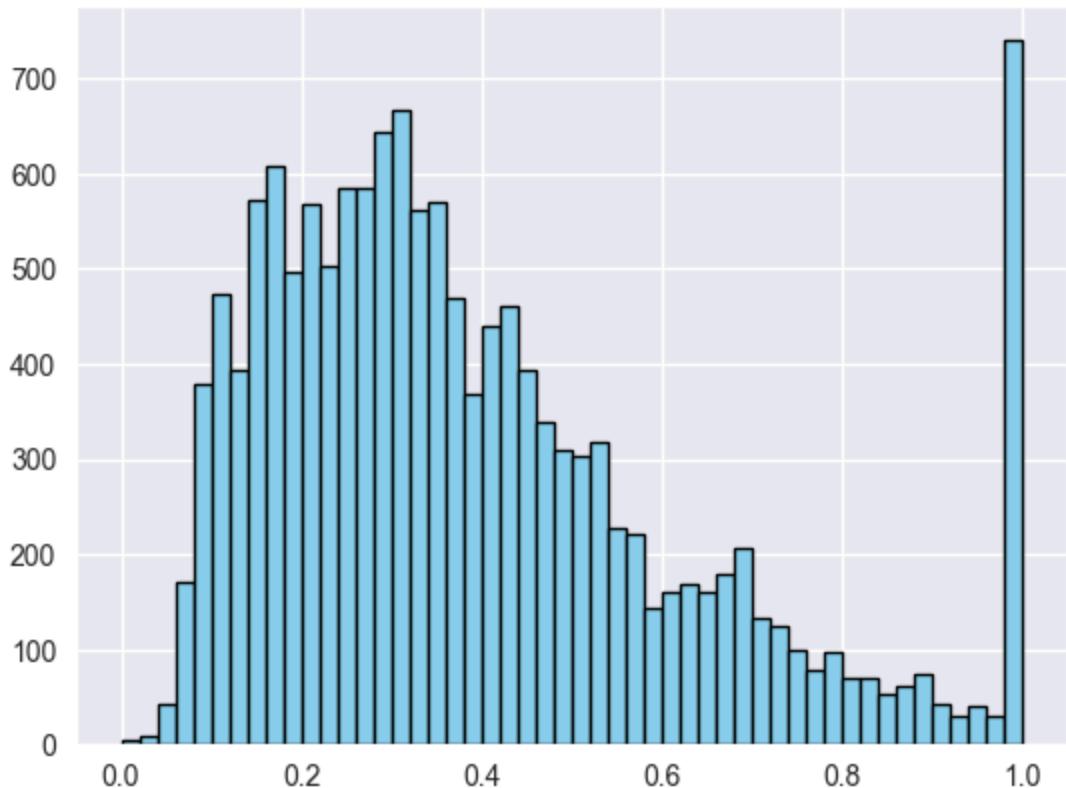
```
In [37]: quantile = preprocessing.QuantileTransformer(output_distribution='normal', random_state=42)
        .fit_transform(y_train)
test_transformations(y_train, [(quantile, "Quantile")])
```



Sqrt transformation was applied to the predicted variable. Because it spans a large range of values a MinMax scaler should then be used to further normalize it

```
In [38]: house_value_transformer = Pipeline([
    # ('sqrt', preprocessing.FunctionTransformer(np.sqrt, validate=True)),
    ('scaler', preprocessing.MinMaxScaler())
])

y_train_s = house_value_transformer.fit_transform(y_train)
_ = plt.hist(y_train_s, bins=50, color='skyblue', edgecolor='black')
```



Feature selection

```
In [39]: x_train.var()
```

```
Out[39]: median_income      3.628676e+00
housing_median_age     1.591038e+02
total_rooms            4.678804e+06
total_bedrooms         1.762209e+05
population             1.300113e+06
households              1.460931e+05
latitude                4.561395e+00
longitude               4.011729e+00
is_san_francisco       1.548901e-01
is_los_angeles          2.354123e-01
dtype: float64
```

All features have sufficiently high variance

```
In [40]: selector = SelectKBest(f_regression, k='all')
selector.fit_transform(x_train, y_train)
scores = selector.scores_

ranking = (
    pd.DataFrame({
        "Feature": x_train.columns,
        "F_score": scores,
    })
    .sort_values(by="F_score", ascending=False)
    .reset_index(drop=True)
)

print(ranking)
```

	Feature	F_score
0	median_income	13000.094763
1	is_san_francisco	1469.736783
2	is_los_angeles	1049.902026
3	latitude	295.272032
4	total_rooms	273.956091
5	housing_median_age	165.884984
6	households	63.002647
7	total_bedrooms	37.099764
8	longitude	35.263954
9	population	8.546627

```
C:\Users\matze\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\utils\validation.py:1339: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
```

As was noted before features "total_rooms", "total_bedrooms", "population" and "households" have significant correlation, since population has the weakest correlation with

predicted variable we can discard it. Also "total_bedrooms" is going to be removed, since it is less correlated with predicted variable then "total_rooms"

```
In [41]: def save_transformed(transformer, col_order: List[str], x_train, x_test, x_eval, y_-
    assert col_order[-1] == "house_value"
    data = x_train.copy()
    data["house_value"] = y_train
    preprocessor.fit(data)
    X_train = preprocessor.transform(data)
    X_train = pd.DataFrame(
        X_train,
        columns=col_order
    )

    data = x_test.copy()
    data["house_value"] = y_test
    X_test = preprocessor.transform(data)
    X_test = pd.DataFrame(
        X_test,
        columns=col_order
    )

    data = x_val.copy()
    data["house_value"] = y_val
    X_val = preprocessor.transform(data)
    X_val = pd.DataFrame(
        X_val,
        columns=col_order
    )

    os.makedirs(path, exist_ok=True)
    X_train.to_csv(f"{path}/train.csv", index=False)
    X_test.to_csv(f"{path}/test.csv", index=False)
    X_val.to_csv(f"{path}/eval.csv", index=False)

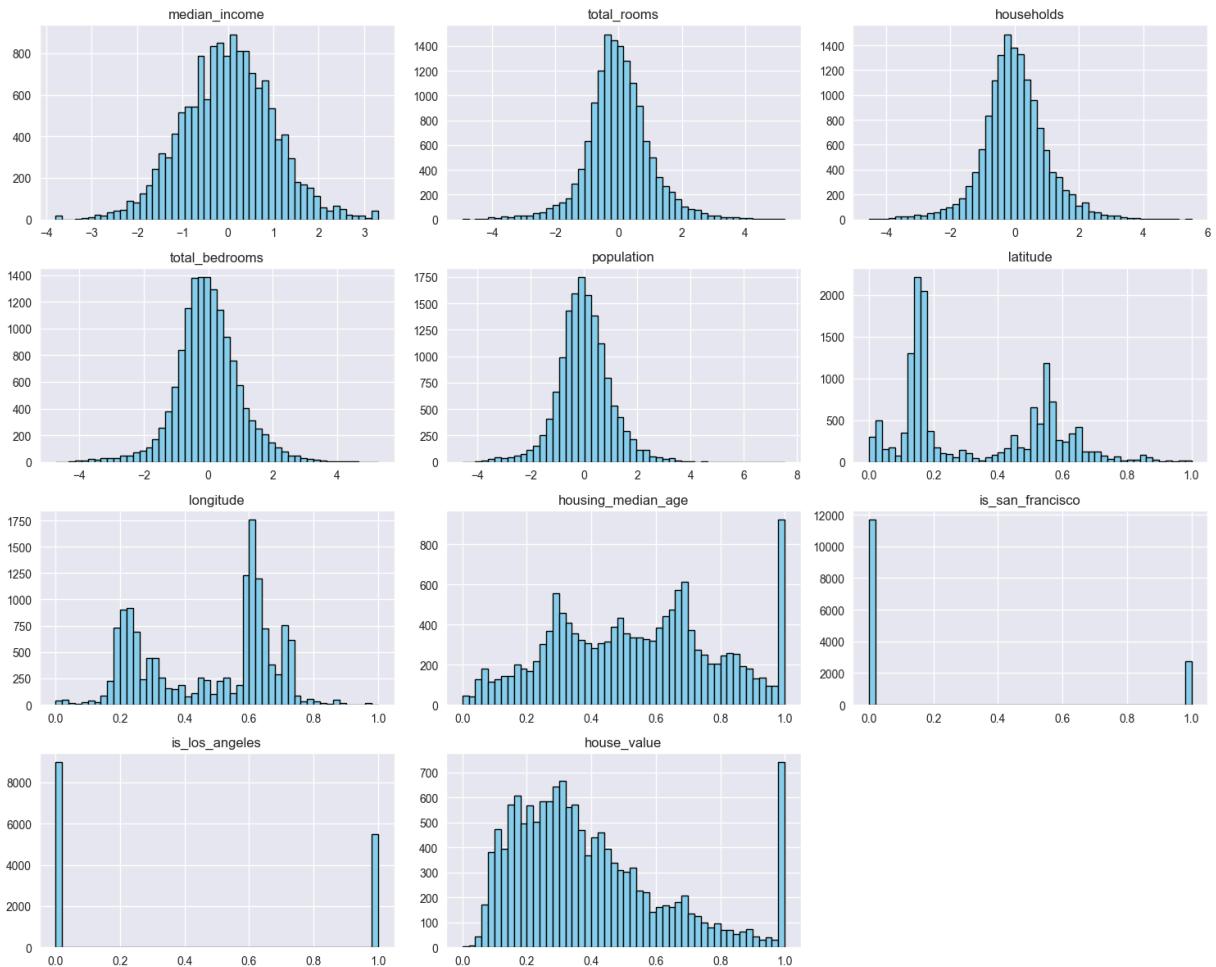
    fitted_house_value_transformer = transformer.named_transformers_["predicted"]
    joblib.dump(fitted_house_value_transformer, f"{path}/house_value_scaler.pkl")
    return X_train, X_test, X_val
```

```
In [42]: from sklearn.compose import ColumnTransformer

preprocessor = ColumnTransformer(
    transformers=[
        ("num", preprocessing.PowerTransformer(method='box-cox'), ["median_income",
            "households", "total_bedrooms", "population"]),
        ("loc", preprocessing.MinMaxScaler(), ["latitude", "longitude", "housing_median_age"]),
        ("bool", preprocessing.FunctionTransformer(), ["is_san_francisco", "is_los_angeles"]),
        ("predicted", house_value_transformer, ["house_value"])
    ]
)
col_order = (
    ["median_income", "total_rooms", "households", "total_bedrooms", "population"]
    ["latitude", "longitude", "housing_median_age"] +
    ["is_san_francisco", "is_los_angeles"] +
    ["house_value"]
```

```
)  
path = f'{TRANSFORMED_PATH}/full_features'  
X_train, _, _ = save_transformed(preprocessor, col_order, x_train, x_test, x_val, y
```

```
In [43]: X_train.hist(bins=50, figsize=(15, 12), color='skyblue', edgecolor='black')  
plt.tight_layout()  
plt.show()
```



```
preprocessor = sklearn.compose.ColumnTransformer(  
    transformers=[  
        ("num", preprocessing.PowerTransformer(method='box-cox'), ["median_income",  
        ("loc", preprocessing.MinMaxScaler(), ["housing_median_age"]),  
        ("bool", preprocessing.FunctionTransformer(), ["is_san_francisco", "is_los_  
        ("predicted", house_value_transformer, ["house_value"])  
    ]  
)  
    col_order = (  
        ["median_income", "total_rooms"] +  
        ["housing_median_age"] +  
        ["is_san_francisco", "is_los_angeles"] +  
        ["house_value"]  
    )  
    path = f'{TRANSFORMED_PATH}/small'  
    X_train, _, _ = save_transformed(preprocessor, col_order, x_train, x_test, x_val, y
```

```
In [45]: X_train.hist(bins=50, figsize=(15, 12), color='skyblue', edgecolor='black')
plt.tight_layout()
plt.show()
```

