

Import libraries

```
In [1]: import itertools
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
import pandas as pd
from typing import List
import wandb
import joblib
import matplotlib.pyplot as plt
import numpy as np
import json
import os
import csv
```

1. Configuration

1.1 Init

```
In [2]: WANDB_PROJECT_NAME = "zneus-project-1"
dataset_used = "full_features"
path = f"data/transformed/{dataset_used}" #path of used dataset
IS_WANDB = True
run_id = 1

#using seed for consistent experiments
RANDOM_SEED = 42
np.random.seed(RANDOM_SEED)
torch.manual_seed(RANDOM_SEED)

#if cuda available, use it
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Device used: {device}")

#model parameters
EPOCH = 50
BATCH_SIZE = 128
LR = 0.001

#transformer for inverse transforms
target_transformer = joblib.load(f"{path}/house_value_scaler.pkl")
```

Device used: cuda

1.1.1 Wandb init

```
In [3]: def init_wandb(name=f"{dataset_used}_{run_id}", config_dict=None):
    if IS_WANDB:
```

```

    default_config = {
        "batch_size": BATCH_SIZE,
        "epoch": EPOCH,
        "lr": LR,
        "loss_fn": "MSELoss",
        "dataset_path": path,
        "activation": "LeakyReLU",
        "optimizer": "Adam",
        "weight_decay": 0.0
    }

    #use dict from tuning
    if config_dict is not None:
        default_config.update(config_dict)

    wandb.init(
        project=f"{WANDB_PROJECT_NAME}",
        name=name,
        config=default_config
)

```

1.2 Model parameters

In [4]:

```

class CSVDataset(Dataset):
    def __init__(self, path: str):
        df = pd.read_csv(path)
        self.X = torch.tensor(df.iloc[:, :-1].values, dtype=torch.float32)
        self.y = torch.tensor(df.iloc[:, -1].values, dtype=torch.float32).view(-1, 1)

    def __len__(self):
        return len(self.X)

    def __getitem__(self, idx):
        return self.X[idx], self.y[idx]

#load datasets from disk
def get_datasets(path: str) -> List[CSVDataset]:
    return [
        CSVDataset(f"{path}/train.csv"),
        CSVDataset(f"{path}/test.csv"),
        CSVDataset(f"{path}/eval.csv")
]

```

1.3 Model class - variants

1.3.1 Base

In [5]:

```

class SimpleNN(nn.Module):
    def __init__(self, input_size: int):
        super(SimpleNN, self).__init__()
        self.net = nn.Sequential(
            nn.Linear(input_size, 128),
            nn.LeakyReLU(0.1),

```

```
        nn.Linear(128, 64),
        nn.LeakyReLU(0.1),
        nn.Linear(64, 32),
        nn.LeakyReLU(0.1),
        nn.Linear(32, 1)
    )

    def forward(self, x):
        return self.net(x)
```

1.3.2 Dropout

```
In [6]: class SimpleNN_dropout(nn.Module):
    def __init__(self, input_size: int):
        super(SimpleNN_dropout, self).__init__()
        self.net = nn.Sequential(
            nn.Linear(input_size, 128),
            nn.LeakyReLU(0.1),
            nn.Dropout(0.1),
            nn.Linear(128, 64),
            nn.LeakyReLU(0.1),
            nn.Dropout(0.3),
            nn.Linear(64, 32),
            nn.LeakyReLU(0.1),
            nn.Linear(32, 1)
        )

    def forward(self, x):
        return self.net(x)
```

1.3.3 Batch Normalization

```
In [7]: class SimpleNN_batch_norm(nn.Module):
    def __init__(self, input_size: int):
        super(SimpleNN_batch_norm, self).__init__()
        self.net = nn.Sequential(
            nn.Linear(input_size, 128),
            nn.BatchNorm1d(128),
            nn.LeakyReLU(0.3),

            nn.Linear(128, 64),
            nn.BatchNorm1d(64),
            nn.LeakyReLU(0.1),

            nn.Linear(64, 32),
            nn.BatchNorm1d(32),
            nn.LeakyReLU(0.1),

            nn.Linear(32, 1)
        )

    def forward(self, x):
        return self.net(x)
```

1.3.4 Skip connections

```
In [8]: class SimpleNN_Res(nn.Module): #skip connections
    def __init__(self, input_size: int):
        super(SimpleNN_Res, self).__init__()
        self.fc1 = nn.Linear(input_size, 128)
        self.act1 = nn.LeakyReLU(0.1)

        self.fc2 = nn.Linear(128, 64)
        self.act2 = nn.LeakyReLU(0.1)

        self.fc3 = nn.Linear(64, 32)
        self.act3 = nn.LeakyReLU(0.1)

        self.out = nn.Linear(32, 1)

    #projection for skip connections
    self.skip_proj = nn.Linear(input_size, 64)

    def forward(self, x):
        #first layer
        x1 = self.act1(self.fc1(x))

        #skip connection: input x + projection to 64
        skip = self.skip_proj(x)
        x2 = self.act2(self.fc2(x1) + skip) #residual connection

        #next layer
        x3 = self.act3(self.fc3(x2))

        out = self.out(x3)
        return out
```

1.3.5 Bottleneck layers

```
In [9]: class SimpleNN_Bottleneck(nn.Module):
    def __init__(self, input_size: int, bottleneck_size: int = 16):
        super(SimpleNN_Bottleneck, self).__init__()

        self.fc1 = nn.Linear(input_size, 128)
        self.act1 = nn.LeakyReLU(0.3)

        #bottleneck layer
        self.fc2 = nn.Linear(128, bottleneck_size)
        self.act2 = nn.LeakyReLU(0.1)

        #output layer
        self.fc3 = nn.Linear(bottleneck_size, 32)
        self.act3 = nn.LeakyReLU(0.1)

        self.out = nn.Linear(32, 1)

    def forward(self, x):
        x1 = self.act1(self.fc1(x))
```

```

        x2 = self.act2(self.fc2(x1))
        x3 = self.act3(self.fc3(x2))
        out = self.out(x3)
    return out

```

1.4 Load data

```
In [10]: #Load datasets and create Loss function
train_df, test_df, eval_df = get_datasets(path)
loss_fn = nn.MSELoss()

def create_loaders(batch_size):
    #divide datasets into batches
    train_loader = DataLoader(train_df, batch_size=batch_size, shuffle=True, worker_
    test_loader = DataLoader(test_df, batch_size=batch_size, shuffle=True, worker_i_
    eval_loader = DataLoader(eval_df, batch_size=batch_size, shuffle=True, worker_i

    return train_loader, test_loader, eval_loader

train_loader, test_loader, eval_loader = create_loaders(BATCH_SIZE) #for base exper
```

1.5 Inverse transform init

```
In [11]: #inverse transform the input for original values
def get_inverse_transformed(y, y_transformer=None):
    if y_transformer is not None:
        y = y.cpu().numpy().reshape(-1, 1)
        y = y_transformer.inverse_transform(y)
    return torch.tensor(y, dtype=torch.float32, device=device)
return y
```

1.6 Evaluate model

```
In [22]: def evaluate(eloader: DataLoader, model, loss_fn, y_transformer=None, is_test=False
model.eval()
total_test_loss = 0
total_test_loss_original = 0
num_of_batches = len(eloader)

mse_losses, rmse_losses = [], [] #for test only

with torch.no_grad():
    if is_test:
        p = "TRAINING" if is_train else "TESTING"
        print(f"\n====={p}=====")

    for index, (X, y) in enumerate(eloader):
        X, y = X.to(device), y.to(device)
        y_pred = model(X)

        loss_transformed = loss_fn(y_pred, y).item()
```

```

        mse_losses.append(loss_transformed)
        total_test_loss += loss_transformed

        y_pred = get_inverse_transformed(y_pred, y_transformer)
        y = get_inverse_transformed(y, y_transformer)

        loss_original = loss_fn(y_pred, y).item()
        total_test_loss_original += loss_original
        if is_test:
            rmse = loss_original ** 0.5
            rmse_losses.append(rmse)

    #print(f"{p.lower().capitalize()} eval: Batch {index + 1}: MSE={total_test_loss/num_of_batches}, total_test_loss_original/{num_of_batches},

    return total_test_loss/num_of_batches, total_test_loss_original/num_of_batches,

```

1.7 Train model

```

In [21]: def train(train: CSVDataset, train_loader: DataLoader, eval_loader: DataLoader, loss_fn):
    model = model if model else SimpleNN_Bottleneck(train.X.shape[1]).to(device) #if None else model

    optimizer = optimizer if optimizer else optim.Adam(model.parameters(), lr=LR) #
    train_mse, eval_mse = [], []
    train_rmse, eval_rmse = [], []

    if not tuning: print("\n====TRAINING====")
    for epoch in range(EPOCH):
        model.train()
        total_train_mse = 0
        total_train_rmse = 0

        for X, y in train_loader:
            X, y = X.to(device), y.to(device)
            optimizer.zero_grad()

            #predict y
            y_pred = model(X)

            #calculate loss
            loss = loss_fn(y_pred, y)
            total_train_mse += loss.item()

            #backward pass
            loss.backward()
            optimizer.step()

        #calculate non-normalized loss, for logging
        with torch.no_grad():
            #inverse transform predicted and true y values
            y_orig = get_inverse_transformed(y.detach(), transformer)
            y_pred_orig = get_inverse_transformed(y_pred.detach(), transformer)

            #calculate mse and rmse from results
            mse = loss_fn(y_pred_orig, y_orig).item()
            rmse = mse ** 0.5

```

```

        total_train_rmse += rmse

    #average train mse/rmse per epoch
    avg_train_mse = total_train_mse / len(train_loader)
    avg_train_rmse = total_train_rmse / len(train_loader)

    #test on eval set
    eval_mse_e, eval_mse_original, _, _, _ = evaluate(eval_loader, model, loss_
    eval_rmse_e = eval_mse_original ** 0.5

    #add to arrays
    train_mse.append(avg_train_mse)
    eval_mse.append(eval_mse_e)
    train_rmse.append(avg_train_rmse)
    eval_rmse.append(eval_rmse_e)

    #if not tuning:
    #print(f"Epoch {epoch+1:03d}: train RMSE={avg_train_rmse:.4f}, eval RMS

return model, train_mse, eval_mse, train_rmse, eval_rmse

```

1.8 Plotting the results locally

```

In [14]: def plot_res(tr_mse, ev_mse, tr_rmse, ev_rmse, te_mse_losses, te_rmse_losses, test_
#plot mse losses
plt.figure(figsize=(8, 5))
plt.plot(range(1, len(tr_mse) + 1), tr_mse, label="Training MSE", linewidth=2)
plt.plot(range(1, len(ev_mse) + 1), ev_mse, label="Evaluation MSE", linewidth=2)
plt.title("Training, Evaluation MSE Over Epochs")
plt.xlabel("Epoch")
plt.ylabel("MSE")
plt.legend()
plt.grid(True, linestyle="--", alpha=0.6)
plt.tight_layout()
plt.show()

#plot rmse - real losses in €
plt.figure(figsize=(8, 5))
plt.plot(range(1, len(tr_rmse) + 1), tr_rmse, label="Training RMSE", linewidth=2)
plt.plot(range(1, len(ev_rmse) + 1), ev_rmse, label="Evaluation RMSE", linewidth=2)
plt.title("Training, Evaluation RMSE Over Epochs")
plt.xlabel("Epoch")
plt.ylabel("RMSE (€)")
plt.legend()
plt.grid(True, linestyle="--", alpha=0.6)
plt.tight_layout()
plt.show()

#test plots - mse
plt.figure(figsize=(8, 5))
plt.plot(range(1, test_batches + 1), te_mse_losses, label="Test MSE", linewidth=2)
plt.title("Test MSE Over Batches")
plt.xlabel("Epoch")
plt.ylabel("MSE")

```

```

plt.legend()
plt.grid(True, linestyle="--", alpha=0.6)
plt.tight_layout()
plt.show()

#test plots - rmse
plt.figure(figsize=(8, 5))
plt.plot(range(1, test_batches + 1), te_rmse_losses, label="Test RMSE", linewidth=2)
plt.title("Test RMSE Over Batches")
plt.xlabel("Epoch")
plt.ylabel("RMSE (€)")
plt.legend()
plt.grid(True, linestyle="--", alpha=0.6)
plt.tight_layout()
plt.show()

```

1.9 Train the model and log the results into WandB and plot locally if allowed in init

```

In [15]: def run_experiment(model=None, optimizer=None, train_loader=None, test_loader=None,
#start wandb run
    if tuning:
        init_wandb(tuning_param[0], tuning_param[1])
    else:
        init_wandb()

#train
model, tr_mse, ev_mse, tr_rmse, ev_rmse = train(train_df, train_loader, eval_loader)

#evaluate on train dataset, seen data
train_mse, train_mse_original, _, _, _ = evaluate(train_loader, model, loss_fn,
train_eval_rmse = train_mse_original ** 0.5

#evaluate test dataset, unseen data
test_mse, test_mse_original, te_mse_losses, te_rmse_losses, test_batches = evaluate(test_loader, model, loss_fn,
te_rmse = test_mse_original ** 0.5

#log results
if IS_WANDB:
    for epoch, (mse_t, mse_e, rmse_t, rmse_e) in enumerate(zip(tr_mse, ev_mse,
        wandb.log({
            "epoch": epoch,
            "train_MSE": mse_t,
            "eval_MSE": mse_e,
            "train_RMSE": rmse_t,
            "eval_RMSE": rmse_e,
        })
    )

    for batch_idx, (mse, rmse) in enumerate(zip(te_mse_losses, te_rmse_losses),
        wandb.log({
            "batch": batch_idx,
            "test_MSE": mse,
            "test_RMSE": rmse
        })
)

```

```

    wandb.finish()

    print(f"Train RMSE: {train_eval_rmse:.4f}, Test RMSE: {te_rmse:.4f}")

    if not tuning: #during tuning we dont plot graphs
        plot_res(tr_mse, ev_mse, tr_rmse, ev_rmse, te_mse_losses, te_rmse_losses, t)

    return train_eval_rmse, te_rmse

```

Code may output different Test RMSE, that is because we are logging batch RMSE to wandb and then wandb calculates a mean, this is not accurate for comparing models. We calculate RMSE on the whole dataset, that is the correct approach.

```
In [23]: run_experiment(train_loader=train_loader, test_loader=test_loader, eval_loader=eval
run_id += 1
```

Tracking run with wandb version 0.22.3

Run data is saved locally in C:\zeus\projekt1\wandb\run-20251109_205457-odofb419

Syncing run [full_features_1](#) to [Weights & Biases](#) (docs)

View project at <https://wandb.ai/xhanushchak-slovensk-technick-univerzita-v-bratislave/zneus-project-1>


View run at <https://wandb.ai/xhanushchak-slovensk-technick-univerzita-v-bratislave/zneus-project-1/runs/odofb419>

=====TRAINING=====

=====TRAINING=====

=====TESTING=====

Run history:



Run summary:

| | |
|------------|-------------|
| batch | 49 |
| epoch | 50 |
| eval_MSE | 0.01189 |
| eval_RMSE | 52894.9021 |
| test_MSE | 0.00943 |
| test_RMSE | 47086.17598 |
| train_MSE | 0.01185 |
| train_RMSE | 52123.38831 |

View run **full_features_1** at: <https://wandb.ai/xhanushchak-slovensk-technick-univerzita-v-bratislave/zneus-project-1/runs/odofb4l9>


View project at: <https://wandb.ai/xhanushchak-slovensk-technick-univerzita-v-bratislave/zneus-project-1>

Synced 4 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)


Find logs at: .\wandb\run-20251109_205457-odofb4l9\logs


Train RMSE: 52090.0550, Test RMSE: 55346.3453

Training, Evaluation MSE Over Epochs



Training, Evaluation RMSE Over Epochs





1.10 Hyperparameter tuning

```
In [17]: def hyperparameter_tuning(train_set: CSVDataset):
    run_id = 0
    best_rmse = float("inf")
    best_config = None
    csv_file = "tuning_results.csv"

    #create file if it does not exist
```

```

if not os.path.exists(csv_file):
    with open(csv_file, mode="w", newline="") as f:
        writer = csv.writer(f)
        writer.writerow([
            "iteration", "activation", "optimizer", "lr",
            "weight_decay", "batch_size", "train_RMSE", "test_RMSE", "new_best"
        ])

#parameter sets
activation_options = [nn.ReLU, lambda: nn.LeakyReLU(0.1)]
optimizer_options = ['Adam', 'SGD']
learning_rates = {
    'Adam': [1e-4, 3e-4, 1e-3],
    'SGD': [1e-2, 3e-3, 1e-3]
}
batch_sizes = [64, 128]
weight_decays = [0, 1e-4]

param_grid = list(itertools.product(
    activation_options,
    optimizer_options,
    batch_sizes,
    weight_decays
))

total_runs = sum(len(learning_rates[o]) for _, o, _, _ in param_grid)
print(f"Starting hyperparameter tuning: {total_runs} total runs\n")

#for effectivity, lets init all needed data loaders needed for experiments
loaders = []
for b in batch_sizes:
    train_loader, test_loader, eval_loader = create_loaders(b)
    loaders.append([train_loader, test_loader, eval_loader]) #index of b in bat

for act_fn, opt_name, batch, wd in param_grid:
    loader_idx = batch_sizes.index(batch)
    for lr in learning_rates[opt_name]:
        activation = act_fn()
        model = SimpleNN_Res(train_set.X.shape[1]).to(device)
        model.act1 = activation
        model.act2 = activation
        model.act3 = activation

        #optimizer
        if opt_name == 'Adam':
            optimizer = optim.Adam(model.parameters(), lr=lr, weight_decay=wd)
        elif opt_name == 'SGD':
            optimizer = optim.SGD(model.parameters(), lr=lr, momentum=0.9, weig

exp_name = f"{dataset_used}_parameter_tuning_{run_id}"
config_dict = {
    "activation": act_fn.__name__ if hasattr(act_fn, "__name__") else "LeakyReLU",
    "batch_size": batch,
    "epoch": EPOCH,
    "lr": lr,
}

```

```

        "loss_fn": "MSELoss",
        "dataset_path": path,
        "optimizer": opt_name,
        "weight_decay": wd
    }

    print(f"[{run_id+1}/{total_runs}] Starting: "
          f"{config_dict['optimizer']} | act={config_dict['activation']} | "
          f"lr={lr} | wd={wd} | batch={batch}")

    train_rmse, test_rmse = run_experiment(model=model, optimizer=optimizer)

    new_best = False #if it the new best model yet
    if test_rmse < best_rmse:
        best_rmse = test_rmse
        best_config = config_dict
        new_best = True
        print(f"New best RMSE: {best_rmse:.4f} with config {best_config}")

    #write results to csv
    with open(csv_file, mode="a", newline="") as f:
        writer = csv.writer(f)
        writer.writerow([run_id+1, config_dict["activation"], opt_name, lr,
                        train_rmse, test_rmse])

    run_id += 1

    if best_config is not None:
        with open("best_config.json", "w") as f:
            json.dump({"best_rmse": best_rmse, "config": best_config}, f, indent=4)
        print(f"\nBest configuration saved to best_config.json:")
        print(json.dumps(best_config, indent=4))
    else:
        print("No valid configuration found!")

```

In [18]: `#hyperparameter_tuning(train_df) #start hyperparameter tuning`

In [20]: `#best iterations`
`tuning_res = pd.read_csv("models_tuning/tuning_improvements.csv")`
`tuning_res`

| | iteration | activation | optimizer | lr | weight_decay | batch_size | train_RMSE | t |
|---|-----------|----------------|-----------|--------|--------------|------------|--------------|-----|
| 0 | 1 | ReLU | Adam | 0.0001 | 0.0 | 64 | 53373.203790 | 568 |
| 1 | 2 | ReLU | Adam | 0.0003 | 0.0 | 64 | 51067.636097 | 556 |
| 2 | 3 | ReLU | Adam | 0.0010 | 0.0 | 64 | 48283.271055 | 552 |
| 3 | 9 | ReLU | Adam | 0.0010 | 0.0 | 128 | 49679.079022 | 549 |
| 4 | 27 | LeakyReLU(0.1) | Adam | 0.0010 | 0.0 | 64 | 47995.374228 | 539 |

1.11 Best Model Results

```
In [24]: with open("models_tuning/best_config.json", "r") as f: #Load best config
    best = json.load(f)

    best_rmse = best["best_rmse"]
    config = best["config"]
    print(f"Best RMSE: {best_rmse}")
    print("Config:", config)

    #create model
    model = SimpleNN_Res(train_df.X.shape[1]).to(device)

    #create activation function
    activation = nn.LeakyReLU(0.1)

    #apply activation
    model.act1 = activation
    model.act2 = activation
    model.act3 = activation

    #create optimizer
    lr = config["lr"]
    weight_decay = config["weight_decay"]

    if config["optimizer"] == "Adam":
        optimizer = optim.Adam(model.parameters(), lr=lr, weight_decay=weight_decay)

    #create loader
    batch_size = config["batch_size"]
    train_loader, test_loader, eval_loader = create_loaders(batch_size)

    #run experiment
    train_rmse, test_rmse = run_experiment(
        model=model,
        optimizer=optimizer,
        train_loader=train_loader,
        test_loader=test_loader,
        eval_loader=eval_loader,
        tuning=False #print results also
    )
```

Best RMSE: 53968.44942545371
 Config: {'activation': 'LeakyReLU(0.1)', 'batch_size': 64, 'epoch': 50, 'lr': 0.001, 'loss_fn': 'MSELoss', 'dataset_path': 'data/transformed/full_features', 'optimizer': 'Adam', 'weight_decay': 0}

Tracking run with wandb version 0.22.3

Run data is saved locally in C:\zeus\projekt1\wandb\run-20251109_205618-gvlsqjyq

Syncing run [full_features_1](#) to Weights & Biases (docs)

View project at <https://wandb.ai/xhanushchak-slovensk-technick-univerzita-v-bratislave/zneus-project-1>


View run at <https://wandb.ai/xhanushchak-slovensk-technick-univerzita-v-bratislave/zneus-project-1/runs/gvlsqjyq>

====TRAINING====

====TRAINING====

====TESTING====

Run history:



Run summary:

| | |
|------------|-------------|
| batch | 49 |
| epoch | 50 |
| eval_MSE | 0.01135 |
| eval_RMSE | 51665.07952 |
| test_MSE | 0.00616 |
| test_RMSE | 38054.22321 |
| train_MSE | 0.01059 |
| train_RMSE | 49376.0815 |


View run **full_features_1** at: <https://wandb.ai/xhanushchak-slovensk-technick-univerzita-v-bratislave/zneus-project-1/runs/gvlsqjyq>

View project at: <https://wandb.ai/xhanushchak-slovensk-technick-univerzita-v-bratislave/zneus-project-1>


Synced 4 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: .\wandb\run-20251109_205618-gvlsqjyq\logs

Train RMSE: 48883.2118, Test RMSE: 54498.1586



Test MSE Over Batches



Test RMSE Over Batches

